

Overview of Java Parallel Streams: Phases

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

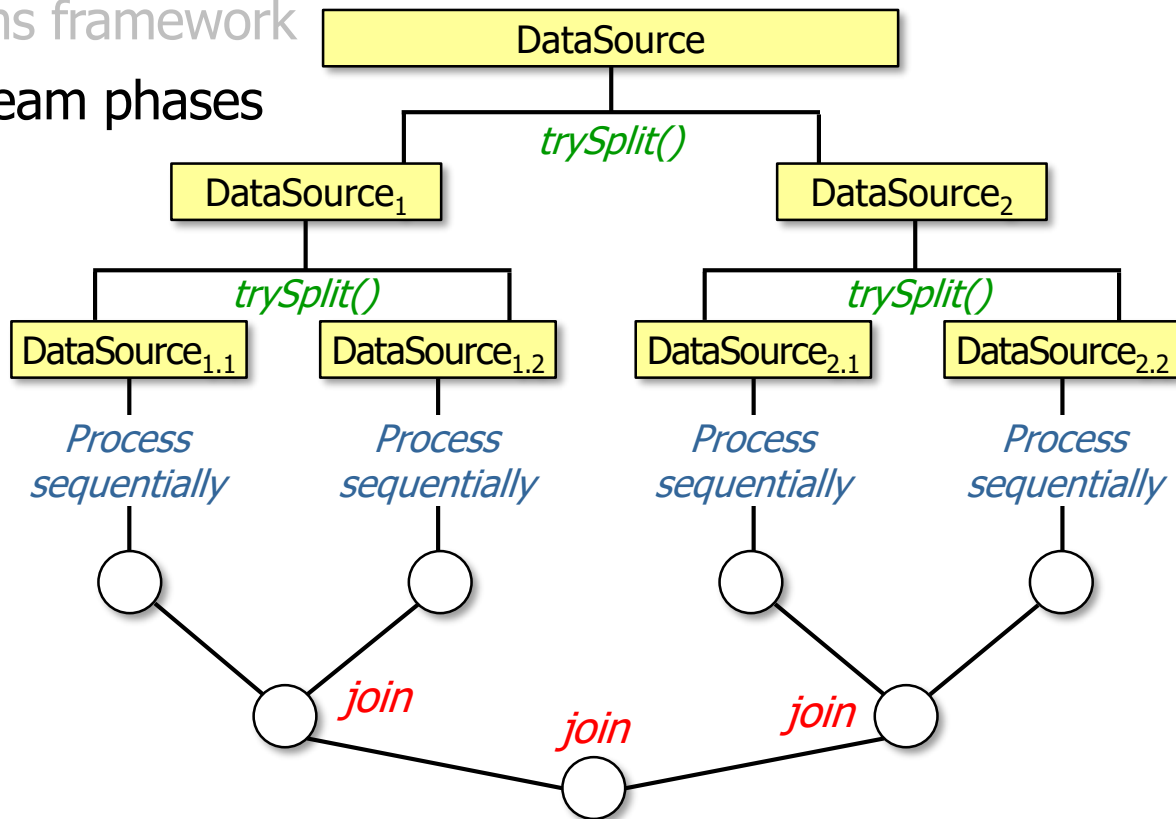
**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Know how aggregate operations & functional programming features are applied in the parallel streams framework
- Be aware of how parallel stream phases work “under the hood”




See www.ibm.com/developerworks/library/j-java-streams-3-brian-goetz

Overview of How a Parallel Stream Works

Overview of How a Parallel Stream Works


- A Java parallel stream implements a “map/reduce” variant optimized for multi-core processors



See en.wikipedia.org/wiki/MapReduce

Overview of How a Parallel Stream Works

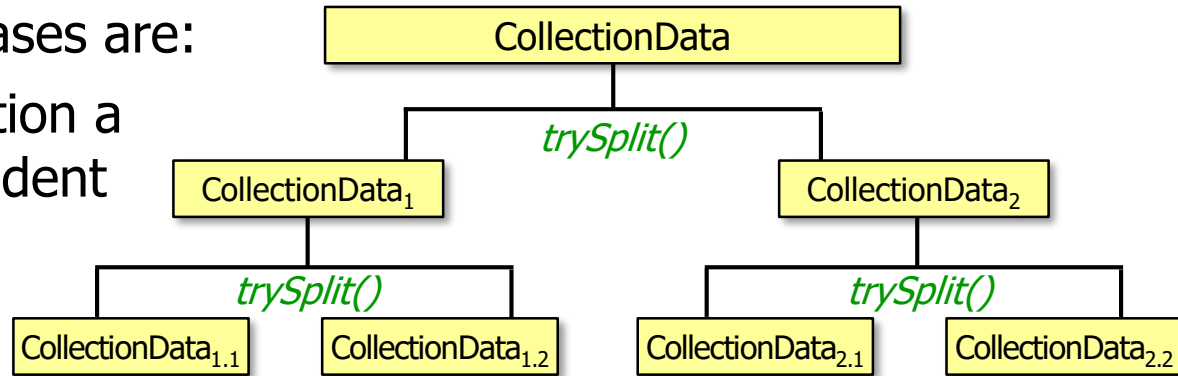
- A Java parallel stream implements a “map/reduce” variant optimized for multi-core processors
- It’s actually a three phase “split-apply-combine” data processing strategy



Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”



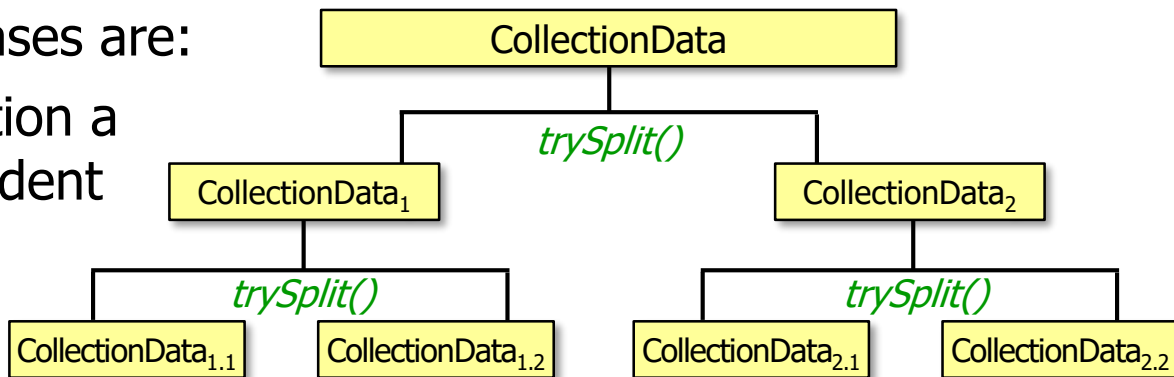
See en.wikipedia.org/wiki/Divide_and_conquer_algorithm

Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

- Spliterators are defined to partition collections in Java



```
public interface Spliterator<T> {  
    boolean tryAdvance(Consumer<? Super T> action);  
  
    Spliterator<T> trySplit();  
  
    long estimateSize();  
  
    int characteristics();  
}
```

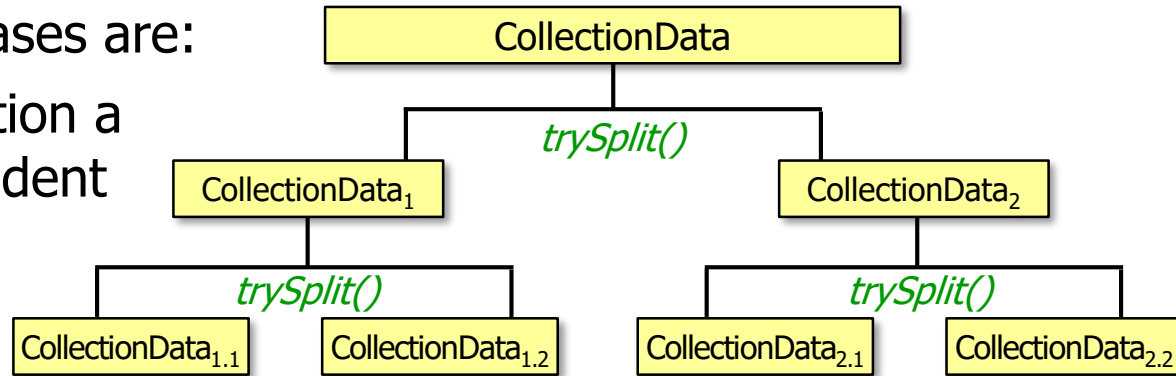
See docs.oracle.com/javase/8/docs/api/java/util/Spliterator.html

Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

- Spliterators are defined to partition collections in Java



*Used for sequential
(& parallel) streams*

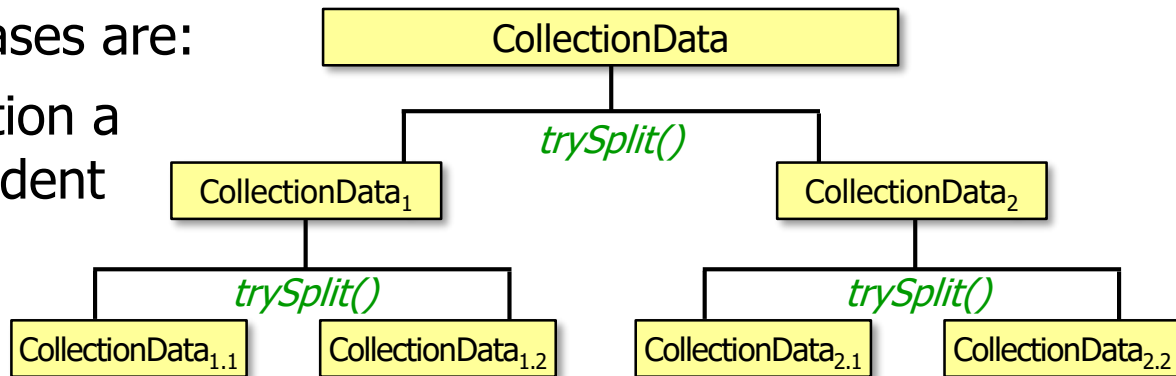
```
public interface Spliterator<T> {  
    boolean tryAdvance(Consumer<? Super T> action);  
  
    Spliterator<T> trySplit();  
  
    long estimateSize();  
  
    int characteristics();  
}
```


Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

- Spliterators are defined to partition collections in Java



*Used only for
parallel streams*

```
public interface Spliterator<T> {  
    boolean tryAdvance(Consumer<? Super T> action);  
  
    Spliterator<T> trySplit();  
  
    long estimateSize();  
  
    int characteristics();  
}
```


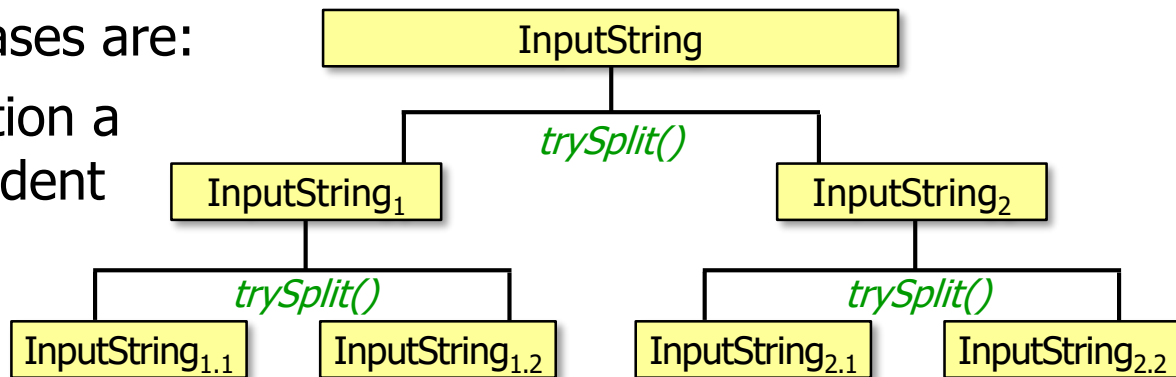
Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

- Splitterators are defined to partition collections in Java

- You can also define custom splitterators

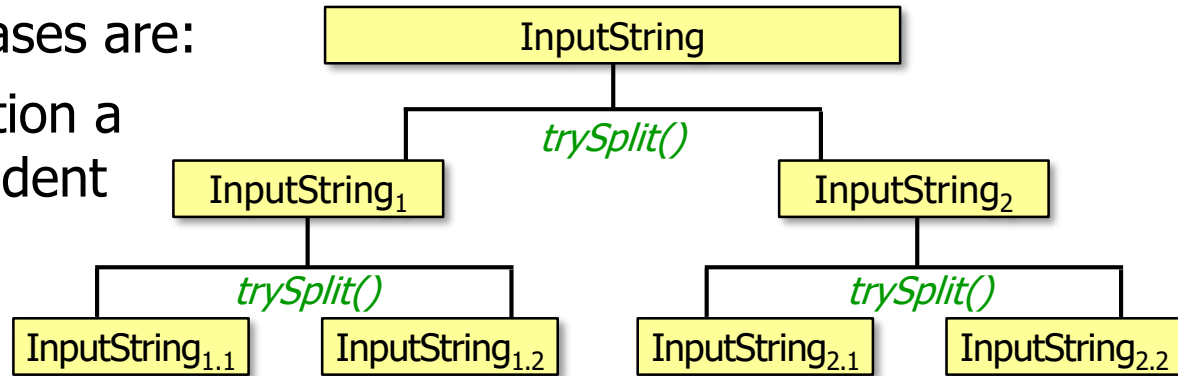


Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

- Splitterators are defined to partition collections in Java
- You can also define custom splitterators
- Parallel streams perform better on data sources that can be split efficiently & evenly

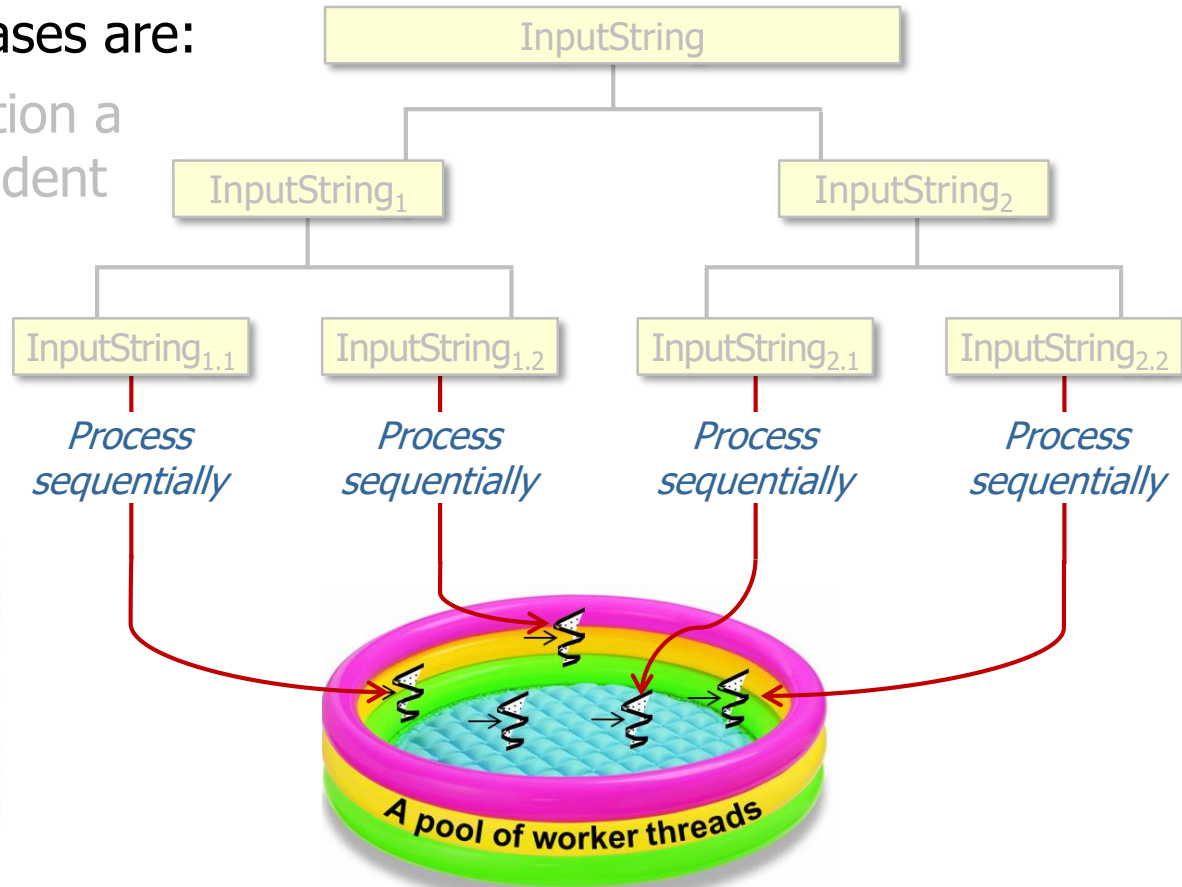


Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. Split – Recursively partition a data source into independent “chunks”

2. Apply – Process chunks independently in one (common) thread pool



Splitting & applying run simultaneously (after certain limit met), not sequentially


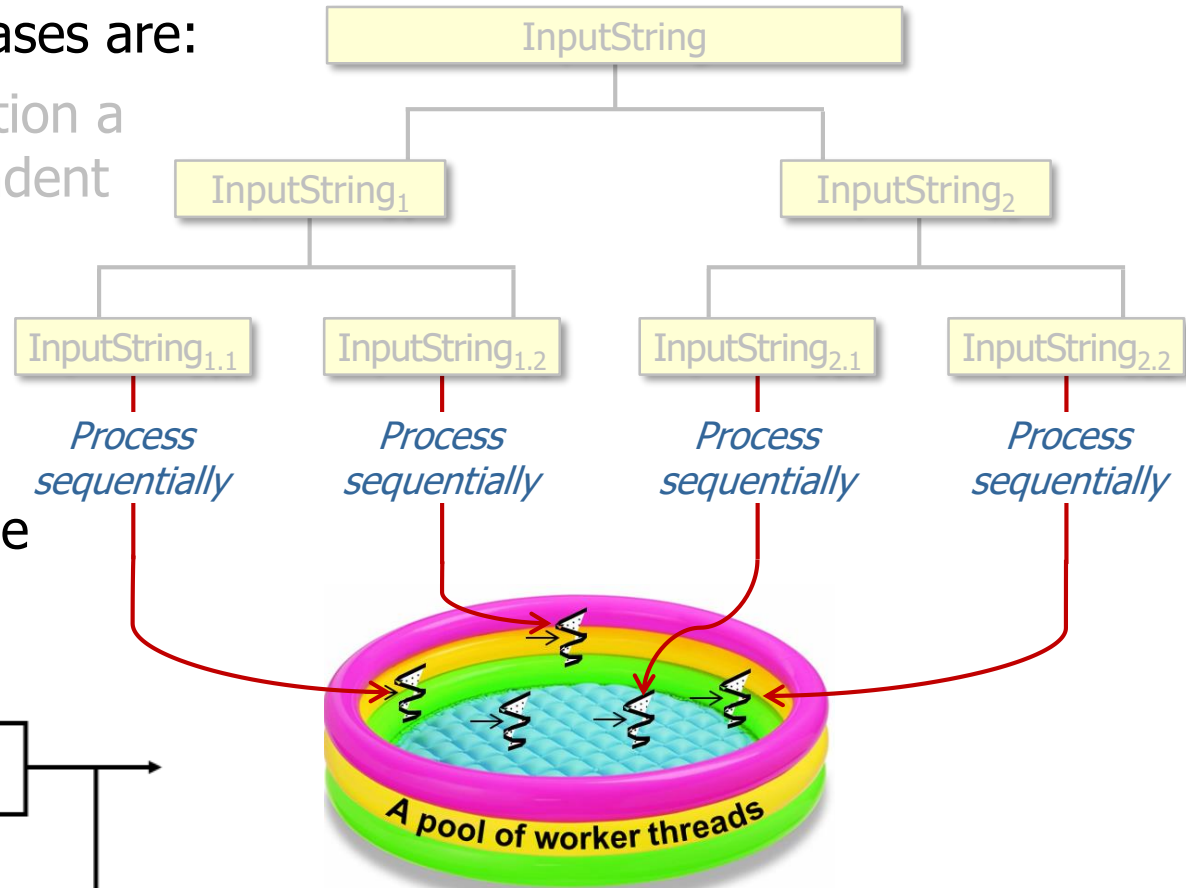
Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. **Split** – Recursively partition a data source into independent “chunks”

2. **Apply** – Process chunks independently in one (common) thread pool

- Programmers have some control over how many threads are in the pool




Overview of How a Parallel Stream Works

- The split-apply-combine phases are:

1. **Split** – Recursively partition a data source into independent “chunks”

2. **Apply** – Process chunks independently in one (common) thread pool

3. **Combine** – Join partial results into a single result



Overview of How a Parallel Stream Works


- The split-apply-combine phases are:

1. **Split** – Recursively partition a data source into independent “chunks”

2. **Apply** – Process chunks independently in one (common) thread pool

3. **Combine** – Join partial results into a single result

- Performed by terminal operations like `collect()` & `reduce()`



End of Overview of Java Parallel Streams: Phases