

Java SearchWithParallelStreams

Example: Evaluating Pros & Cons

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt

Professor of Computer Science

**Institute for Software
Integrated Systems**


**Vanderbilt University
Nashville, Tennessee, USA**





Learning Objectives in this Part of the Lesson

- Know how Java parallel streams are applied in `SearchWithParallelStreams`
- Understand the pros & cons of the `SearchWithParallelStreams` class

<<Java Class>>

 **SearchWithParallelStreams**

 `processStream():List<List<SearchResults>>`

 `processInput(CharSequence):List<SearchResults>`




See [SearchStreamGang/src/main/java/livelessons/streamgangs/SearchWithParallelStreams.java](https://github.com/StreamGang/StreamGang/blob/master/src/main/java/livelessons/streamgangs/SearchWithParallelStreams.java)


Pros of the SearchWith ParallelStreams Class

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!



<<Java Class>>

 **SearchWithParallelStreams**

◆ processStream():List<List<SearchResults>>

■ processInput(CharSequence):List<SearchResults>

See docs.oracle.com/javase/tutorial/collections/streams/parallelism.html

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!

Here's processStream() from SearchWithSequentialStream that we examined earlier

```
List<List<SearchResults>>  
    processStream() {  
        return getInput()  
            .stream()  
            .map(this::processInput)  
            .collect(toList());  
    }
```

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!

```
List<List<SearchResults>>  
    processStream() {  
    return getInput()  
        .stream()  
        .map(this::processInput)  
        .collect(toList());  
}
```

VS


```
List<List<SearchResults>>  
    processStream() {  
    return getInput()  
        .parallelStream()  
        .map(this::processInput)  
        .collect(toList());  
}
```

*Here's processStream() in
SearchWithParallelStreams*

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!

Changing all the stream() calls to parallelStream() calls is the minuscule difference between implementations!!



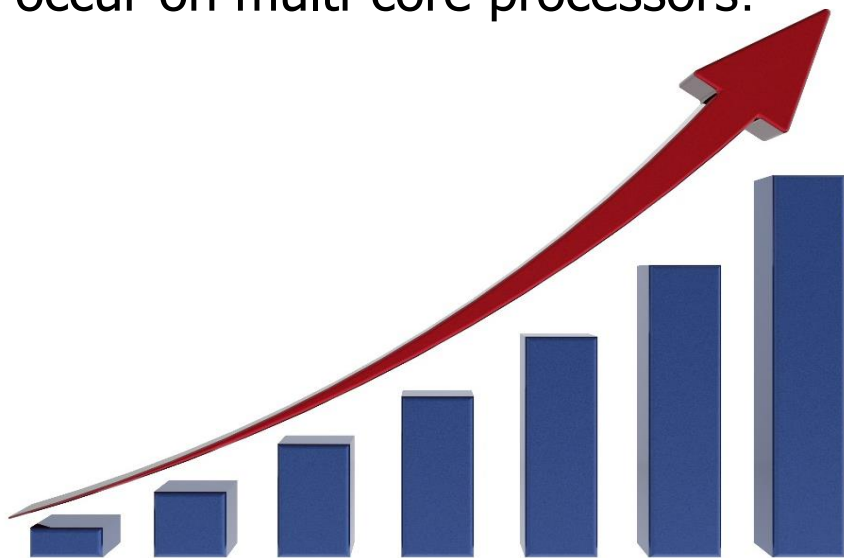
```
List<List<SearchResults>>  
    processStream() {  
        return getInput()  
            .stream()  
            .map(this::processInput)  
            .collect(toList());  
    }
```

VS

```
List<List<SearchResults>>  
    processStream() {  
        return getInput()  
            .parallelStream()  
            .map(this::processInput)  
            .collect(toList());  
    }
```

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!
- Moreover, substantial speedups can occur on multi-core processors!



Input Strings to Search



Search Phrases





```
Starting SearchStreamGangTest
PARALLEL_SPLITATOR executed in 409 msecs
COMPLETABLE_FUTURES_INPUTS executed in 426 msecs
COMPLETABLE_FUTURES_PHASES executed in 427 msecs
PARALLEL_STREAMS executed in 437 msecs
PARALLEL_STREAM_PHASES executed in 440 msecs
RXJAVA_PHASES executed in 485 msecs
PARALLEL_STREAM_INPUTS executed in 802 msecs
RXJAVA_INPUTS executed in 866 msecs
SEQUENTIAL_LOOPS executed in 1638 msecs
SEQUENTIAL_STREAM executed in 1958 msecs
Ending SearchStreamGangTest
```

Tests conducted on a 2.7GHz quad-core Lenovo P50 with 32 Gbytes of RAM

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!
- Moreover, substantial speedups can occur on multi-core processors!

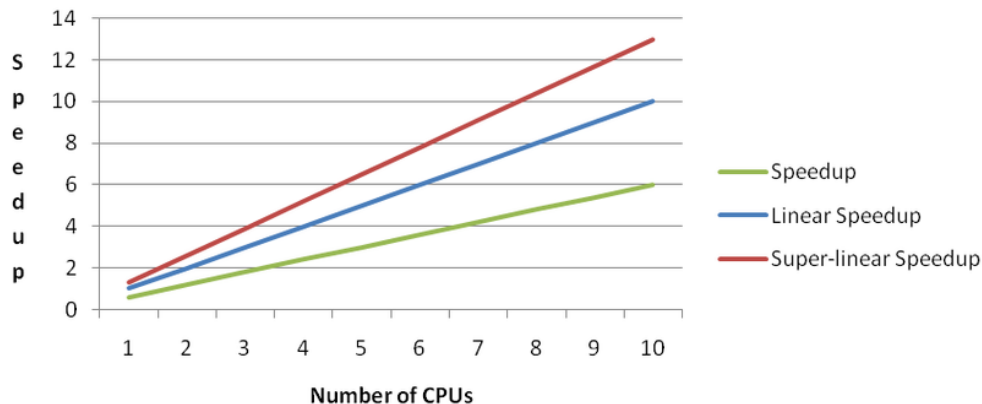


```
Starting SearchStreamGangTest
PARALLEL_SPLITATOR executed in 369 msecs
PARALLEL_STREAMS executed in 373 msecs
COMPLETABLE_FUTURES_INPUTS executed in 377 msecs
COMPLETABLE_FUTURES_PHASES executed in 383 msecs
PARALLEL_STREAM_PHASES executed in 385 msecs
RXJAVA_PHASES executed in 434 msecs
PARALLEL_STREAM_INPUTS executed in 757 msecs
RXJAVA_INPUTS executed in 774 msecs
SEQUENTIAL_LOOPS executed in 1485 msecs
SEQUENTIAL_STREAM executed in 1578 msecs
Ending SearchStreamGangTest
```

Tests conducted on a 2.9GHz quad-core MacBook Pro with 16 Gbytes of RAM

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!
- Moreover, substantial speedups can occur on multi-core processors!
- Superlinear speed-ups arise from “hyper-threaded” (virtual) cores



Input Strings to Search



Search Phrases

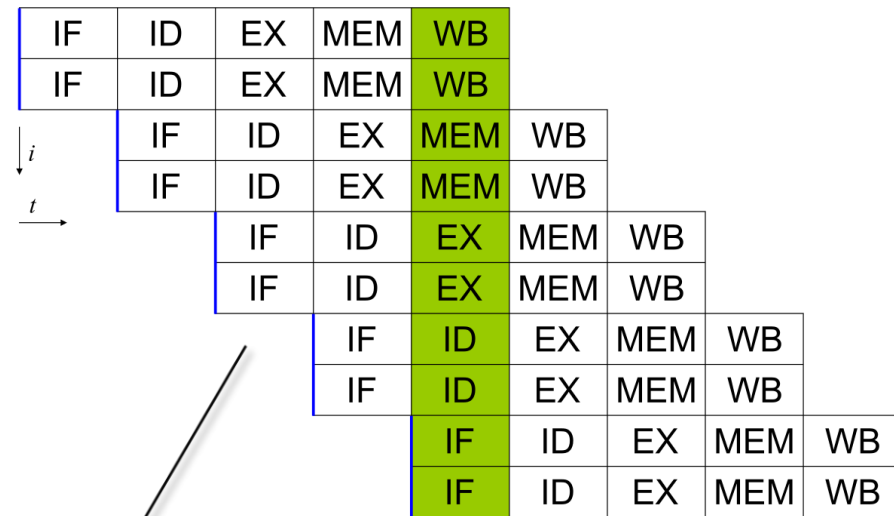


```
Starting SearchStreamGangTest
PARALLEL_SPLITTERATOR executed in 369 msecs
PARALLEL_STREAMS executed in 373 msecs
COMPLETABLE_FUTURES_INPUTS executed in 377 msecs
COMPLETABLE_FUTURES_PHASES executed in 383 msecs
PARALLEL_STREAM_PHASES executed in 385 msecs
RXJAVA_PHASES executed in 434 msecs
PARALLEL_STREAM_INPUTS executed in 757 msecs
RXJAVA_INPUTS executed in 774 msecs
SEQUENTIAL_LOOPS executed in 1485 msecs
SEQUENTIAL_STREAM executed in 1578 msecs
Ending SearchStreamGangTest
```

See en.wikipedia.org/wiki/Hyper-threading

Pros of the SearchWithParallelStreams Class

- This example shows that the difference between sequential & parallel streams is often minuscule!
 - Moreover, substantial speedups can occur on multi-core processors!
- Superlinear speed-ups arise from “hyper-threaded” (virtual) cores
- Increases the # of independent instructions in the pipeline via a superscalar architecture



A superscalar processor can execute more than one instruction during a clock cycle by simultaneously dispatching multiple instructions to different execution units

Cons of the SearchWith ParallelStreams Class

Cons of the SearchWithParallelStreams Class

- Just because two minuscule changes are needed doesn't mean this is the best implementation!

Other Java concurrency/parallelism strategies are even more efficient..

Input Strings to Search



Search Phrases




Starting SearchStreamGangTest

PARALLEL_SPLITATOR executed in 409 msecs
COMPLETABLE_FUTURES_INPUTS executed in 426 msecs
COMPLETABLE_FUTURES_PHASES executed in 427 msecs
PARALLEL_STREAMS executed in 437 msecs
PARALLEL_STREAM_PHASES executed in 440 msecs
RXJAVA_PHASES executed in 485 msecs
PARALLEL_STREAM_INPUTS executed in 802 msecs
RXJAVA_INPUTS executed in 866 msecs
SEQUENTIAL_LOOPS executed in 1638 msecs
SEQUENTIAL_STREAM executed in 1958 msecs
Ending SearchStreamGangTest

Tests conducted on a 2.7GHz quad-core Lenovo P50 with 32 Gbytes of RAM

Cons of the SearchWithParallelStreams Class

- Just because two minuscule changes are needed doesn't mean this is the best implementation!



Input Strings to Search



Search Phrases



```
Starting SearchStreamGangTest
PARALLEL_SPLITATOR executed in 409 msecs
COMPLETABLE_FUTURES_INPUTS executed in 426 msecs
COMPLETABLE_FUTURES_PHASES executed in 427 msecs
PARALLEL_STREAMS executed in 437 msecs
PARALLEL_STREAM_PHASES executed in 440 msecs
RXJAVA_PHASES executed in 485 msecs
PARALLEL_STREAM_INPUTS executed in 802 msecs
RXJAVA_INPUTS executed in 866 msecs
SEQUENTIAL_LOOPS executed in 1638 msecs
SEQUENTIAL_STREAM executed in 1958 msecs
Ending SearchStreamGangTest
```

There's no substitute for systematic benchmarking & experimentation

Cons of the SearchWithParallelStreams Class

- We'll show how to overcome these cons in an upcoming lesson that focuses on the `SearchWithParallelSpliterator` class

```
<<Java Class>>
```

SearchWithParallelSpliterator

- ◆ `processStream():List<List<SearchResults>>`

```
processInput(CharSequence):List<SearchResults>
```

Input Strings to Search



Search Phrases



Starting SearchStreamGangTest

PARALLEL_SPLITTER executed in 409 msec

COMPLETABLE_FUTURES_INPUTS executed in 426 msec

COMPLETABLE_FUTURES_PHASES executed in 427 msecs

PARALLEL_STREAMS executed in 437 msec

PARALLEL_STREAM_PHASES executed in 440 msec

RXJAVA PHASES executed in 485 msec

PARALLEL_STREAM_INPUTS executed in 802 msec

RXJAVA INPUTS executed in 866 msec

SEQUENTIAL LOOPS executed in 1638 msecs

SEQUENTIAL STREAM executed in 1958 msec

Ending SearchStreamGangTest

Cons of the SearchWithParallelStreams Class

- We'll show how to overcome these cons in an upcoming lesson that focuses on the SearchWithParallelSpliterator class

<<Java Class>>


 **SearchWithParallelSpliterator**

 process

 process

ts>>

hResults>



Starting SearchStreamGangTest

PARALLEL_SPLITERATOR executed in 409 msecs

COMPLETABLE_FUTURES_INPUTS executed in 426 msecs

COMPLETABLE_FUTURES_PHASES executed in 427 msecs

PARALLEL_STREAMS executed in 437 msecs

PARALLEL_STREAM_PHASES executed in 440 msecs

RXJAVA_PHASES executed in 485 msecs

PARALLEL_STREAM_INPUTS executed in 802 msecs

RXJAVA_INPUTS executed in 866 msecs

SEQUENTIAL_LOOPS executed in 1638 msecs

SEQUENTIAL_STREAM executed in 1958 msecs

Ending SearchStreamGangTest

SearchWithParallelSpliterator is thus the most aggressively parallelism strategy!

End of Java SearchWith ParallelStreams Example: Evaluating Pros & Cons