

Java Parallel Stream Internals: Configuring the Common Fork-Join Pool

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Understand parallel stream internals, e.g.
 - Know what can change & what can't
 - Partition a data source into "chunks"
 - Process chunks in parallel
- Configure the Java parallel stream common fork-join pool

```
int desiredThreads = 8;  
System.setProperty  
("java.util.concurrent."  
 + "ForkJoinPool.common."  
 + "parallelism",  
desiredThreads);
```



Configuring the Parallel Stream Common Fork-Join Pool

Configuring the Parallel Stream Common Fork-Join Pool

- By default the common ForkJoinPool has one less thread than the # of cores


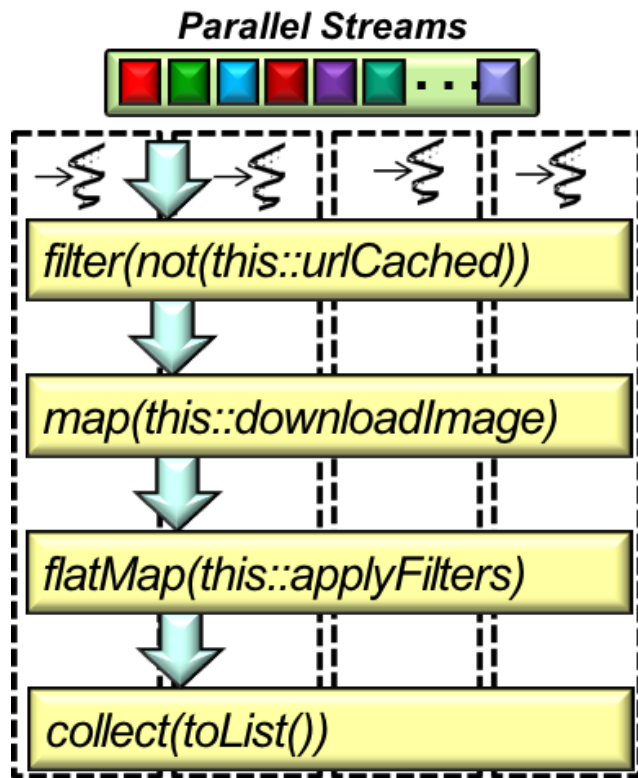
```
System.out.println  
("The parallelism in the"  
+ "common fork-join pool is "  
+ ForkJoinPool  
  .getCommonPoolParallelism());
```

e.g., returns 3 on a quad-core processor



Configuring the Parallel Stream Common Fork-Join Pool



- By default the common ForkJoinPool has one less thread than the # of cores



A parallel stream can use all cores since it uses the invoking thread, e.g., main thread



Configuring the Parallel Stream Common Fork-Join Pool

- However, the default # of fork-join pool threads may be inadequate



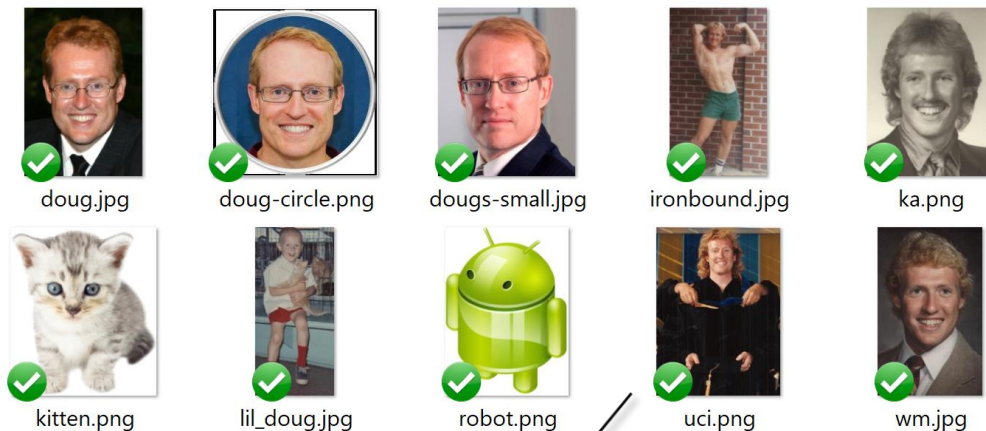
Configuring the Parallel Stream Common Fork-Join Pool

- However, the default # of fork-join pool threads may be inadequate, e.g.
 - Consider a parallel image downloading & processing app




Configuring the Parallel Stream Common Fork-Join Pool

- However, the default # of fork-join pool threads may be inadequate, e.g.
 - Consider a parallel image downloading & processing app



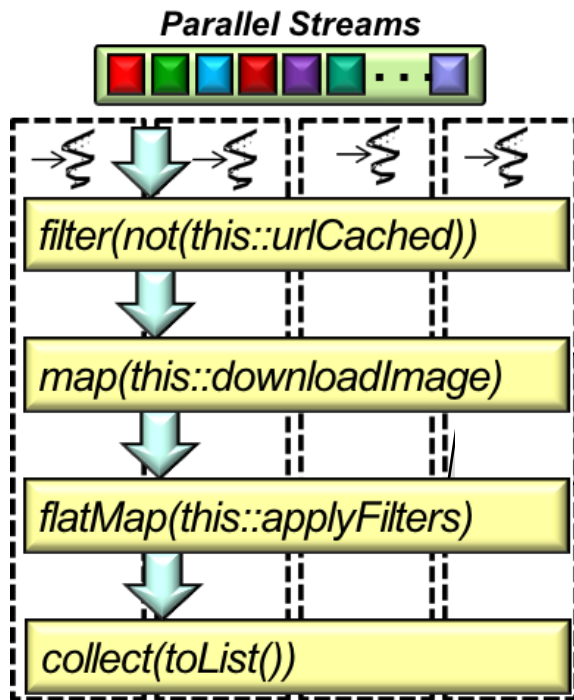
Problems may occur when trying to download more images than # of cores



These problems may range from underutilization of processor cores to deadlock..

Configuring the Parallel Stream Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically




```
int numberOfDownloads = 8;  
System.setProperty  
("java.util.concurrent."  
 + "ForkJoinPool.common."  
 + "parallelism",  
numberOfDownloads);
```



Configuring the Parallel Stream Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically
- Setting this property affects all parallel streams in a process

```
int numberOfDownloads = 8;  
System.setProperty  
("java.util.concurrent."  
 + "ForkJoinPool.common."  
 + "parallelism",  
 numberOfDownloads);
```



It's hard to estimate the total # of threads to set in the common fork-join pool

Configuring the Parallel Stream Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically
 - Setting this property affects all parallel streams in a process
- The ManagedBlocker interface can also be used to add worker threads to common fork-join pool temporarily



```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>  
        (supplier);  
  
...  
ForkJoinPool.managedBlock(mb);  
  
...  
return mb.getResult();
```



Configuring the Parallel Stream Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically
 - Setting this property affects all parallel streams in a process
- The ManagedBlocker interface can also be used to add worker threads to common fork-join pool temporarily
 - This is useful for behaviors that block on I/O and/or synchronizers

```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>  
        (supplier);  
  
...  
ForkJoinPool.managedBlock(mb);  
  
...  
return mb.getResult();
```



Configuring the Parallel Stream Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically
 - Setting this property affects all parallel streams in a process
- The ManagedBlocker interface can also be used to add worker threads to common fork-join pool temporarily
 - This is useful for behaviors that block on I/O and/or synchronizers
 - This interface can only be used with the common fork-join pool..

```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>  
        (supplier);  
...  
ForkJoinPool.managedBlock(mb);  
...  
return mb.getResult();
```



See lessons on "*The Java Fork-Join Pool: the ManagedBlocker Interface*"

End of Java Parallel Stream Internals: Configuring the Common Fork-Join Pool