

# Java Parallel Stream Internals: Combining Results (Part 2)

**Douglas C. Schmidt**

**[d.schmidt@vanderbilt.edu](mailto:d.schmidt@vanderbilt.edu)**

**[www.dre.vanderbilt.edu/~schmidt](http://www.dre.vanderbilt.edu/~schmidt)**



**Professor of Computer Science**

**Institute for Software  
Integrated Systems**

**Vanderbilt University  
Nashville, Tennessee, USA**



# Learning Objectives in this Part of the Lesson

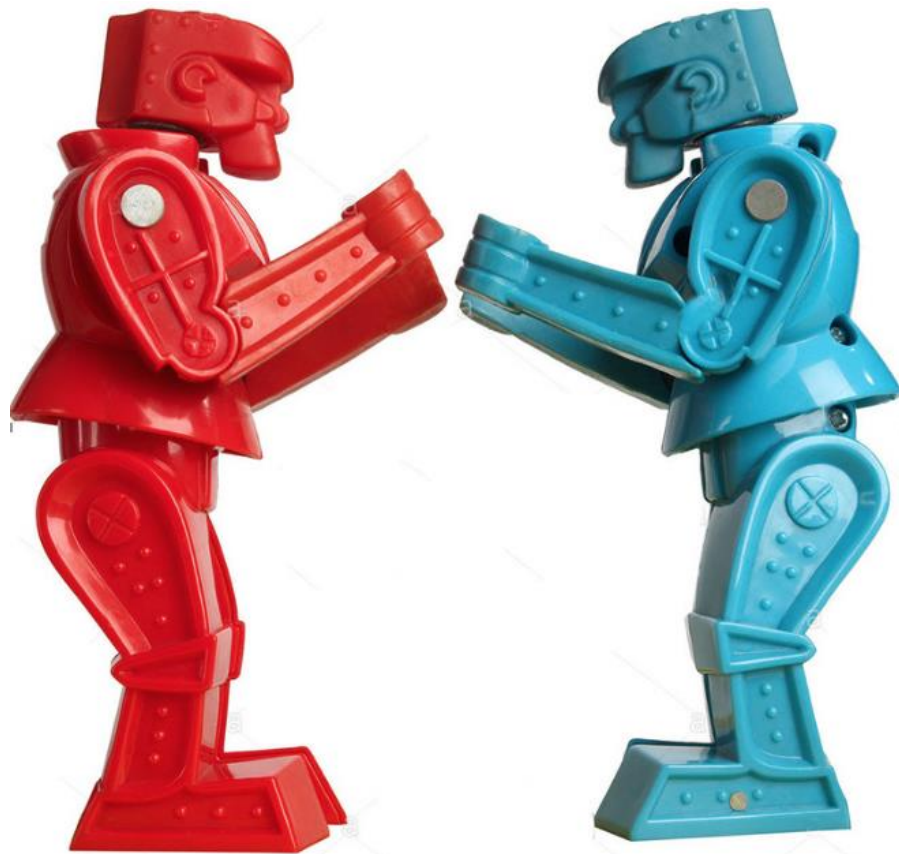
---

- Understand parallel stream internals, e.g.
  - Know what can change & what can't
  - Partition a data source into "chunks"
  - Process chunks in parallel via the common fork-join pool
  - Configure the Java parallel stream common fork-join pool
- Perform a reduction to combine partial results into a single result
  - Be aware of common traps & pitfalls with parallel streams



# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`



# Combining Results in a Parallel Stream

---

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
                StringBuilder::append,
                StringBuilder::append)
        .toString();
}
```

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*Convert a list of words  
into a stream of words*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

Naturally, this call doesn't really do any work since streams are "lazy"

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*A stream can be dynamically switched to "parallel" mode!*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*The "last" call to `.parallel()` or `.sequential()` in a stream "wins"*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```

# Combining Results in a Parallel Stream


- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*This code works when parallel is false since the StringBuilder is only called in a single thread*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```





# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*This code fails when parallel is true since `reduce()` expects to do an "immutable" reduction*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
}
```



# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions



*There's a race condition here since `StringBuilder` is not thread-safe..*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
                StringBuilder::append,
                StringBuilder::append)
        .toString();
}
```

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
  - One solution use `reduce()` with string concatenation

```
void streamReduceConcat
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new String(),
            (x, y) -> x + y);
```

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
  - One solution use `reduce()` with string concatenation

```
void streamReduceConcat
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new String(),
            (x, y) -> x + y);
```

*This simple fix is inefficient due to string concatenation overhead*

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
  - One solution use `reduce()` with string concatenation
  - Another solution uses `collect()` with the joining collector

```
void streamCollectJoining  
    (boolean parallel) {  
    ...  
    Stream<String> wordStream =  
        allWords.stream();  
  
    if (parallel)  
        wordStream.parallel();  
  
    String words = wordStream  
        .collect(joining());  
}
```

# Combining Results in a Parallel Stream

- It's important to understand the semantic differences between `collect()` & `reduce()`, e.g.
- Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
  - One solution use `reduce()` with string concatenation
  - Another solution uses `collect()` with the joining collector

```
void streamCollectJoining  
    (boolean parallel) {  
    ...  
    Stream<String> wordStream =  
        allWords.stream();  
  
    if (parallel)  
        wordStream.parallel();  
  
    String words = wordStream  
        .collect(joining());  
}
```

*This is a much better solution!!*

# Combining Results in a Parallel Stream

- Also beware of issues related to associativity & identity with `reduce()`

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
            // Could use (x, y) -> x + y  
            Math::addExact);  
}
```

# Combining Results in a Parallel Stream

- Also beware of issues related to associativity & identity with `reduce()`

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

*This code fails for a parallel stream since subtraction is not associative*

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
            // Could use (x, y) -> x + y  
            Math::addExact);  
}
```



# Combining Results in a Parallel Stream

- Also beware of issues related to associativity & identity with `reduce()`

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
            // Could use (x, y) -> x + y  
            Math::addExact);  
}
```

*This code fails if  
identity is not 0L*

The “identity” of an OP is defined as “identity OP value == value” (& inverse)

# Combining Results in a Parallel Stream

- Also beware of issues related to associativity & identity with `reduce()`

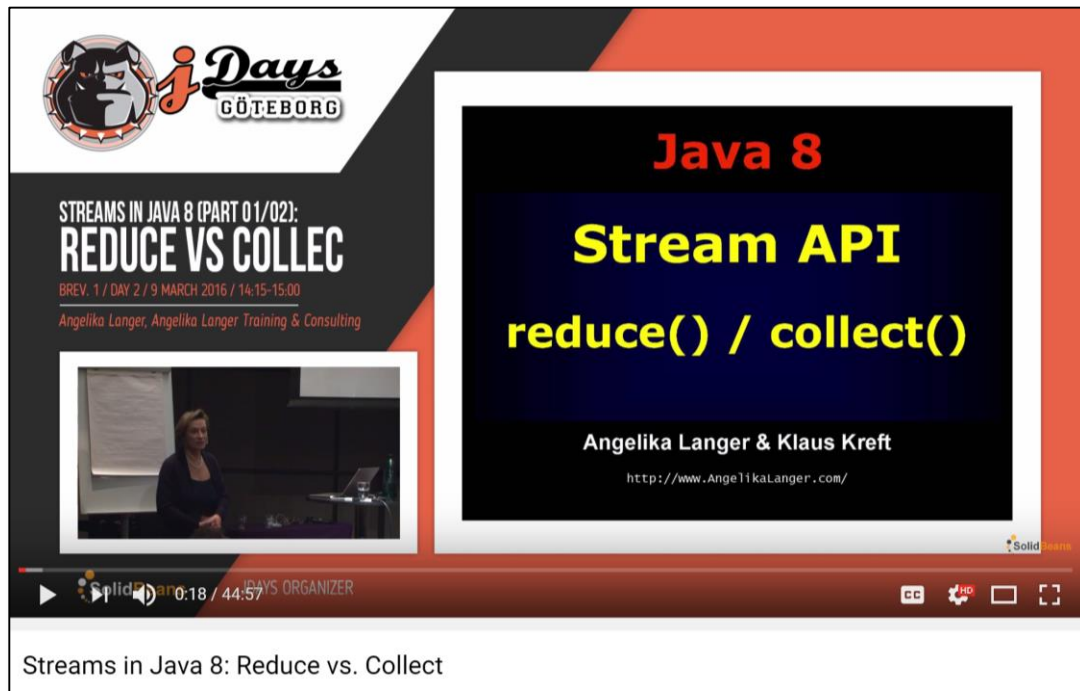
```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
            (x, y) -> x - y);  
}
```

```
void testProd(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
            (x, y) -> x * y);  
}
```

*This code fails if  
identity is not 1L*

# Combining Results in a Parallel Stream

- More good discussions about `reduce()` vs. `collect()` appear online



The image shows a YouTube video player interface. The video title is "STREAMS IN JAVA 8 (PART 01/02): REDUCE VS COLLEC". The presenter is Angelika Langer, Angelika Langer Training & Consulting. The video is part of a series titled "iDays GÖTEBORG". The video content shows a slide titled "Java 8 Stream API reduce() / collect()" by Angelika Langer & Klaus Kreft. The video player controls show the video is at 0:18 / 44:57. The video is licensed under CC BY.

**Streams in Java 8 (Part 01/02): REDUCE VS COLLEC**

BREV. 1 / DAY 2 / 9 MARCH 2016 / 14:15-15:00

Angelika Langer, Angelika Langer Training & Consulting

**Java 8**

**Stream API**

**reduce() / collect()**

Angelika Langer & Klaus Kreft

<http://www.AngelikaLanger.com/>

0:18 / 44:57

Streams in Java 8: Reduce vs. Collect

See [www.youtube.com/watch?v=oWIWEKNM5Aw](http://www.youtube.com/watch?v=oWIWEKNM5Aw)

---

# End of Java Parallel Stream Internals: Combining Results (Part 2)