

The Java Fork-Join Pool: Structure & Functionality (Part 1)

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

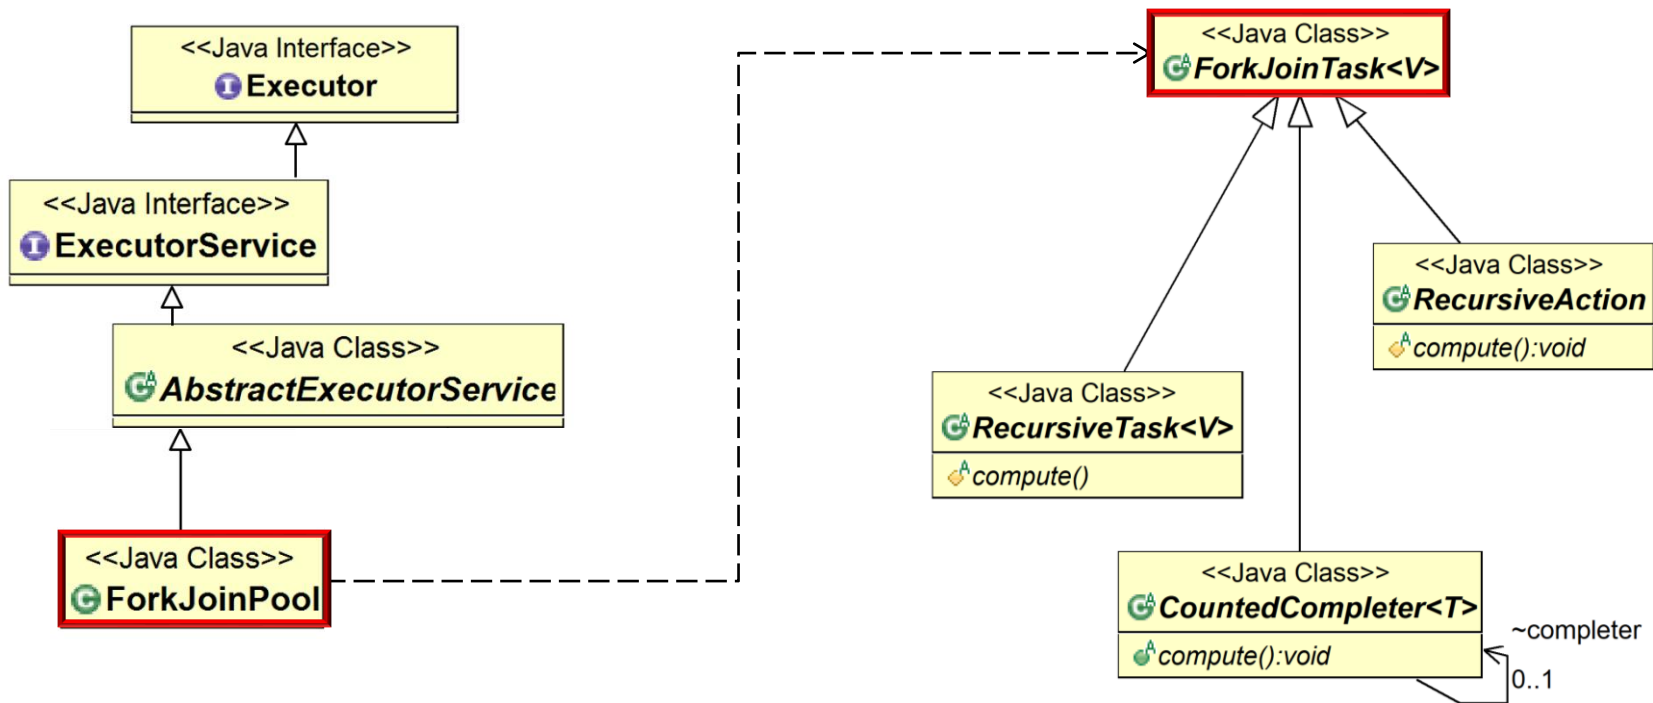
**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Understand how the Java fork-join framework processes tasks in parallel
- Recognize the structure & functionality of the fork-join framework



The Structure & Functionality of the Fork-Join Framework

The Structure & Functionality of the Fork-Join Framework

- ForkJoinPool implements the ExecutorService interface

Class ForkJoinPool

```
java.lang.Object
  java.util.concurrent.AbstractExecutorService
    java.util.concurrent.ForkJoinPool
```

All Implemented Interfaces:

Executor, ExecutorService

```
public class ForkJoinPool
  extends AbstractExecutorService
```

An `ExecutorService` for running `ForkJoinTasks`. A `ForkJoinPool` provides the entry point for submissions from non-`ForkJoinTask` clients, as well as management and monitoring operations.

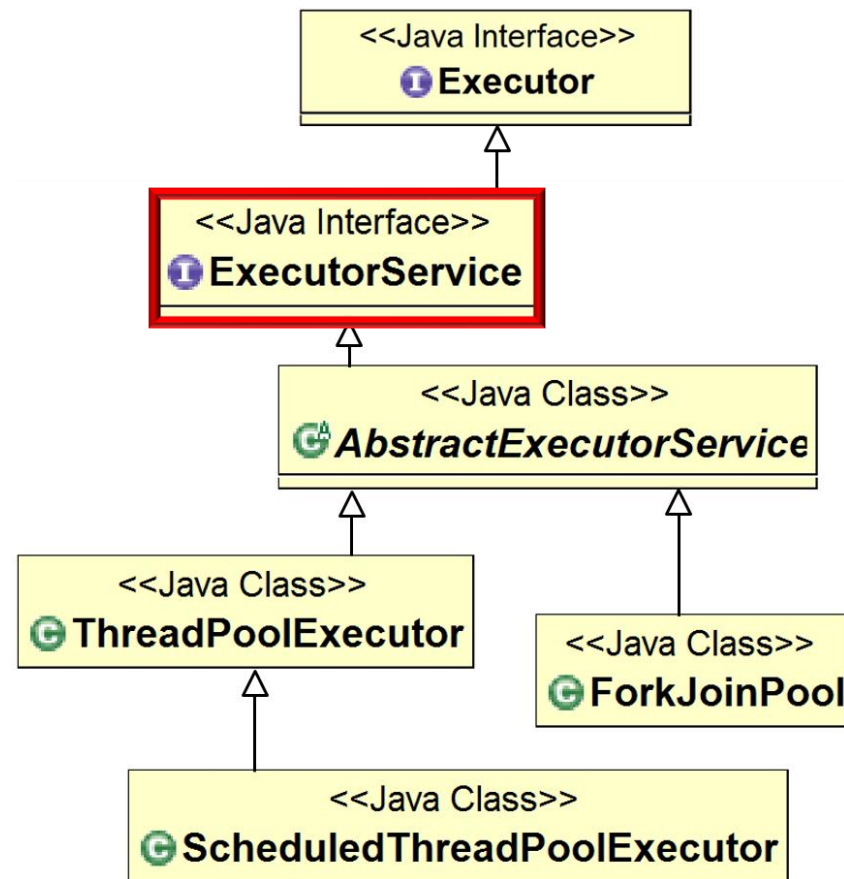
A `ForkJoinPool` differs from other kinds of `ExecutorService` mainly by virtue of employing *work-stealing*: all threads in the pool attempt to find and execute tasks submitted to the pool and/or created by other active tasks (eventually blocking waiting for work if none exist). This enables efficient processing when most tasks spawn other subtasks (as do most `ForkJoinTasks`), as well as when many small tasks are submitted to the pool from external clients. Especially when setting *asyncMode* to true in constructors, `ForkJoinPools` may also be appropriate for use with event-style tasks that are never joined.

A static `commonPool()` is available and appropriate for most applications. The common pool is used by any `ForkJoinTask` that is not explicitly submitted to a specified pool. Using the common pool normally reduces resource usage (its threads are slowly reclaimed during periods of non-use, and reinstated upon subsequent use).

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.html

The Structure & Functionality of the Fork-Join Framework

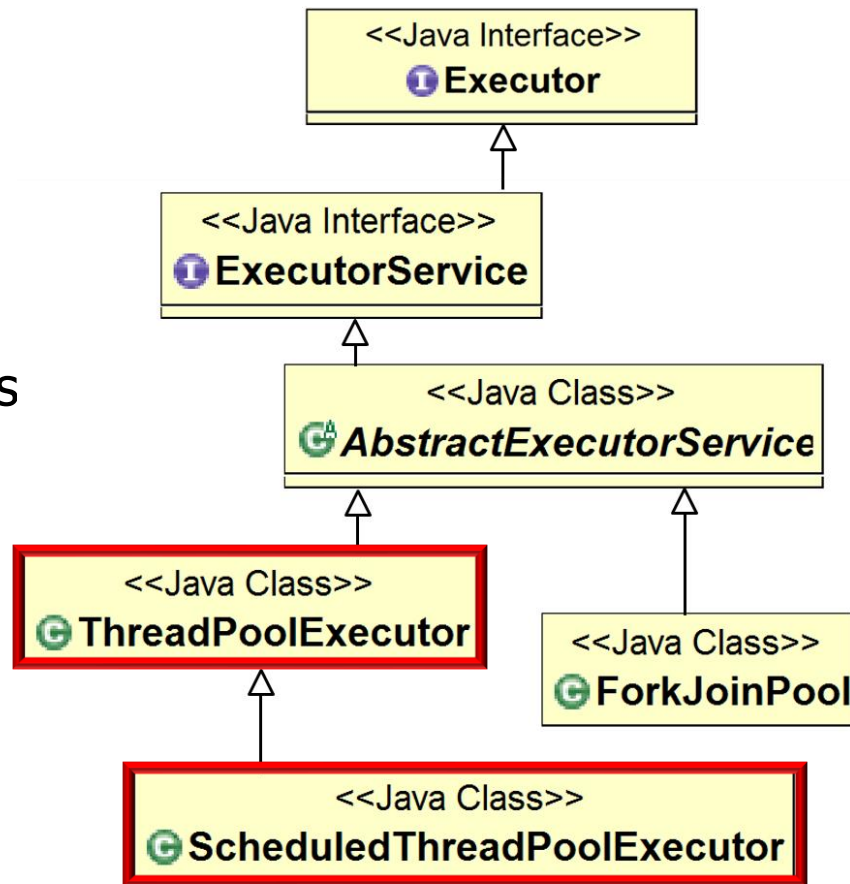


- ForkJoinPool implements the ExecutorService interface
- This interface is the basis for Java Executor framework subclasses



See docs.oracle.com/javase/tutorial/essential/concurrency/executors.html

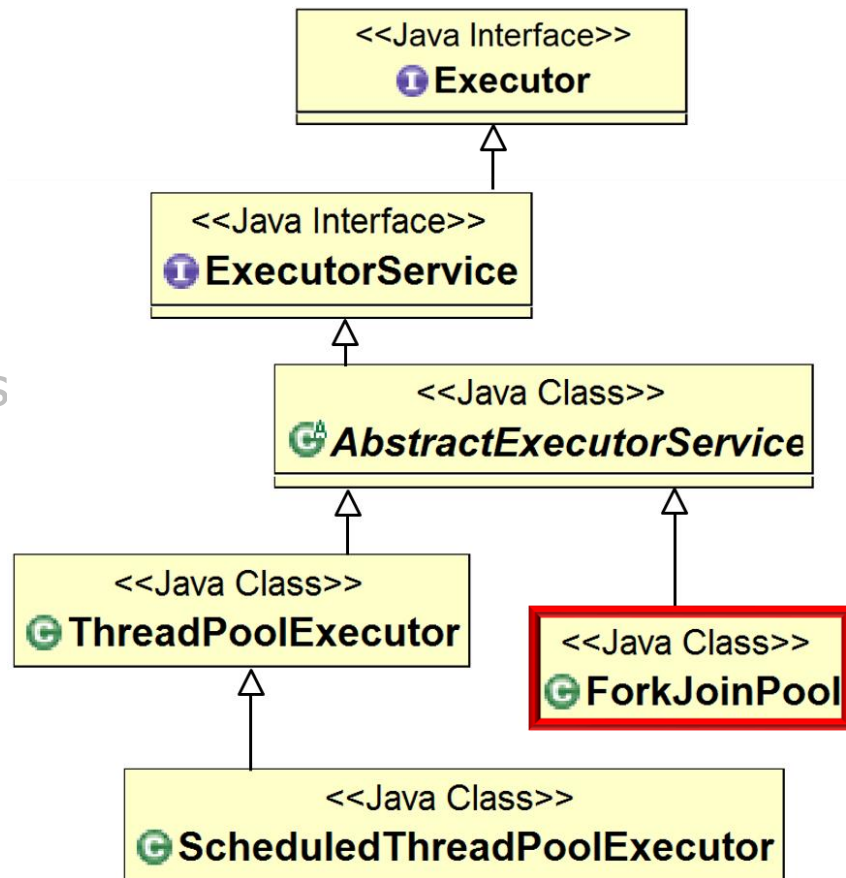
The Structure & Functionality of the Fork-Join Framework

- ForkJoinPool implements the ExecutorService interface
 - This interface is the basis for Java Executor framework subclasses
- Other implementations of Executor Service execute runnables or callables



The Structure & Functionality of the Fork-Join Framework

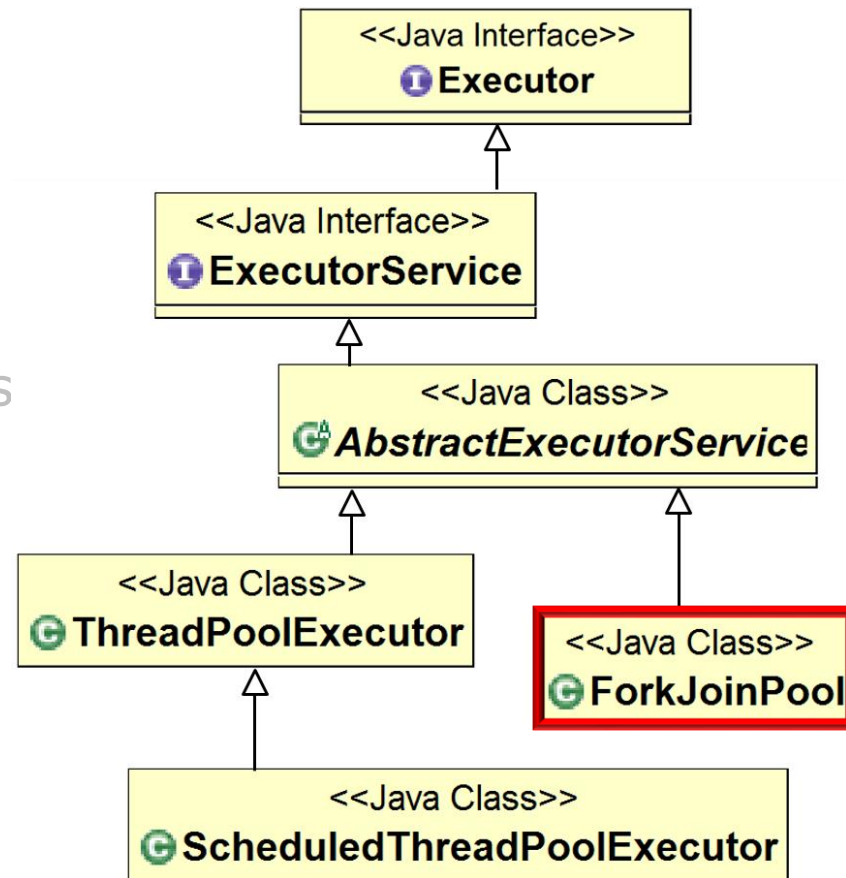
- ForkJoinPool implements the ExecutorService interface
 - This interface is the basis for Java Executor framework subclasses
 - Other implementations of Executor Service execute runnables or callables
- In contrast, the ForkJoinPool executes ForkJoinTasks



The Structure & Functionality of the Fork-Join Framework

- ForkJoinPool implements the ExecutorService interface
 - This interface is the basis for Java Executor framework subclasses
 - Other implementations of Executor Service execute runnables or callables
- In contrast, the ForkJoinPool executes ForkJoinTasks

purpose



It can also execute runnables & callables, but that's not its main purpose

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask associates a chunk of data along with a computation on that data

Class ForkJoinTask<V>

java.lang.Object
java.util.concurrent.ForkJoinTask<V>

All Implemented Interfaces:

Serializable, Future<V>

Direct Known Subclasses:

CountedCompleter, RecursiveAction, RecursiveTask

```
public abstract class ForkJoinTask<V>
    extends Object
    implements Future<V>, Serializable
```

Abstract base class for tasks that run within a `ForkJoinPool`. A `ForkJoinTask` is a thread-like entity that is much lighter weight than a normal thread. Huge numbers of tasks and subtasks may be hosted by a small number of actual threads in a `ForkJoinPool`, at the price of some usage limitations.

A "main" `ForkJoinTask` begins execution when it is explicitly submitted to a `ForkJoinPool`, or, if not already engaged in a ForkJoin computation, commenced in the `ForkJoinPool.commonPool()` via `fork()`, `invoke()`, or related methods. Once started, it will usually in turn start other subtasks. As indicated by the name of this class, many programs using `ForkJoinTask` employ only methods `fork()` and `join()`, or derivatives such as `invokeAll`. However, this class also provides a number of other methods that can come into play in advanced usages, as well as extension mechanics that allow support of new forms of fork/join processing.

See docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinTask.html

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask associates a chunk of data along with a computation on that data
- This enables fine-grained data parallelism



Class ForkJoinTask<V>

java.lang.Object
java.util.concurrent.ForkJoinTask<V>

All Implemented Interfaces:

Serializable, Future<V>

Direct Known Subclasses:

CountedCompleter, RecursiveAction, RecursiveTask


```
public abstract class ForkJoinTask<V>  
extends Object  
implements Future<V>, Serializable
```

Abstract base class for tasks that run within a `ForkJoinPool`. A `ForkJoinTask` is a thread-like entity that is much lighter weight than a normal thread. Huge numbers of tasks and subtasks may be hosted by a small number of actual threads in a `ForkJoinPool`, at the price of some usage limitations.


A "main" `ForkJoinTask` begins execution when it is explicitly submitted to a `ForkJoinPool`, or, if not already engaged in a ForkJoin computation, commenced in the `ForkJoinPool.commonPool()` via `fork()`, `invoke()`, or related methods. Once started, it will usually in turn start other subtasks. As indicated by the name of this class, many programs using `ForkJoinTask` employ only methods `fork()` and `join()`, or derivatives such as `invokeAll`. However, this class also provides a number of other methods that can come into play in advanced usages, as well as extension mechanics that allow support of new forms of fork/join processing.

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask is lighter weight than a Java thread



ForkJoinTask




Thread

e.g., it doesn't maintain its own run-time stack, registers, thread-local storage, etc.


The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask is lighter weight than a Java thread
- A large # of ForkJoinTasks can thus run in a small # of worker threads in a fork-join pool



The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask is lighter weight than a Java thread
- A large # of ForkJoinTasks can thus run in a small # of worker threads in a fork-join pool



Each worker thread has its own stack, registers, etc.




See blog.jamesdbloom.com/JVMInternals.html

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging

Parent ForkJoinTask



Child ForkJoinTasks

ForkJoinTask
<T>


fork() – Arranges to asynchronously execute this task in the appropriate pool

V

join() – Returns result of computation when it is done

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging




ForkJoinTask
<T>

fork() – Arranges to asynchronously execute this task in the appropriate pool

V

join() – Returns result of computation when it is done




ForkJoinTask

fork() is a lightweight variant of Thread.start() that creates a child ForkJoinTask

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging




ForkJoinTask
<T>

fork() – Arranges to asynchronously execute this task in the appropriate pool

V

join() – Returns result of computation when it is done



fork() doesn't run the task, but places it on a work queue in the calling worker thread

The Structure & Functionality of the Fork-Join Framework


- A ForkJoinTask has two methods that control parallel processing/merging

ForkJoinTask
<T>

fork() – Arranges to asynchronously execute this task in the appropriate pool

V


join() – Returns result of computation when it is done



join() returns the result of a child task to the parent task that forked it

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging




ForkJoinTask <T>	fork() – Arranges to asynchronously execute this task in the appropriate pool
V	join() – Returns result of computation when it is done

- Unlike `Thread.join()`, `ForkJoinTask.join()` doesn't simply block the calling thread

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging



ForkJoinTask
<T>

fork() – Arranges to asynchronously execute this task in the appropriate pool

V

join() – Returns result of computation when it is done


- Unlike Thread.join(), ForkJoinTask.join() doesn't simply block the calling thread
 - It uses a worker thread to run tasks



“Pitch in” via the “Collaborative Jiffy Lube” model of processing!

The Structure & Functionality of the Fork-Join Framework

- A ForkJoinTask has two methods that control parallel processing/merging



ForkJoinTask <T>	fork() – Arranges to asynchronously execute this task in the appropriate pool
V	join() – Returns result of computation when it is done

- Unlike Thread.join(), ForkJoinTask.join() doesn't simply block the calling thread
 - It uses a worker thread to run tasks
 - When a worker thread encounters a join() it processes other tasks until it notices the target sub-task is done

End of the Java Fork-Join Pool: Structure & Functionality (Part 1)