

The Java Fork-Join Pool: Implementing `applyAllIter()`

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Apply the fork-join framework in practice
- Examine the applyAllIter() method
 - This method uses “work-stealing” to disperse tasks to worker threads

```
<T> List<T> applyAllIter
(List<T> list,
 Function<T, T> op,
 ForkJoinPool fjPool) {
    return fjPool
        .invoke(new
            RecursiveTask
                <List<T>>() {
                    protected List<T>
                        compute() {
                            ...
                        }
                    }) ;
}
```

Implementing the applyAllIter() Method

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) {
    return fjPool
        .invoke(new RecursiveTask<List<T>>() {
            protected List<T> compute() {
                ...
            }
        });
}
```

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,  
                        ForkJoinPool fjPool) {  
    return fjPool  
        .invoke(new RecursiveTask<List<T>>() {  
            protected List<T> compute() {  
                ...  
            }  
        }) ;  
}
```

These parameters are treated as "effectively final" variables in the anonymous inner class below

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) {
    return fjPool
        .invoke(new RecursiveTask<List<T>>() {
            protected List<T> compute() {
                ...
            }
        });
}
```

Create an anonymous RecursiveTask instance whose compute() method iterative creates many sub-tasks

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) {
    return fjPool
        .invoke(new RecursiveTask<List<T>>() {
            protected List<T> compute() {
                ...
            }
        });
}
```



Code is verbose due to lack of functional interface (& thus can't use lambdas)..

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,  
                          ForkJoinPool fjPool) {
```

```
    return fjPool
```

```
        .invoke(new RecursiveTask<List<T>>() {
```

```
            protected List<T> compute() {
```

```
                ...
```

```
            }
```

```
        });
```

```
    }
```



*Invoke the task on the fork-join pool
& then wait for & return the results*

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```

Hook method implements the main fork-join task

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```

Lists that hold the forked tasks & the results

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```

*Iterate through input list,
fork all the tasks, & add
them to the forks list*

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,  
                        ForkJoinPool fjPool) { ...
```

```
    protected List<T> compute() {
```

```
        List<ForkJoinTask<T>> forks = new LinkedList<>();
```

```
        List<T> res = new LinkedList<>();
```

```
        for (T t : list)
```

```
            forks.add(new RecursiveTask<T>() {
```

```
                protected T compute() { return
```

```
                }.fork());
```

```
        for (ForkJoinTask<T> task : forks) res.add(task.join());
```

```
        return res;
```

```
    } ...
```



This implementation relies on "work-stealing" to disperse tasks to worker threads

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        Join all results of forked tasks & add to results list

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```

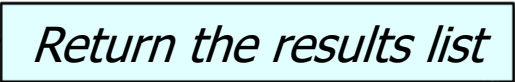
Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```



Return the results list

Implementing the applyAllIter() Method

- Apply an 'op' to all items in the list by calling fork-join methods iteratively

```
<T> List<T> applyAllIter(List<T> list, Function<T, T> op,
                        ForkJoinPool fjPool) { ...
    protected List<T> compute() {
        List<ForkJoinTask<T>> forks = new LinkedList<>();
        List<T> res = new LinkedList<>();

        for (T t : list)
            forks.add(new RecursiveTask<T>() {
                protected T compute() { return op.apply(t); }
            }.fork());

        for (ForkJoinTask<T> task : forks) res.add(task.join());
        return res;
    } ...
```

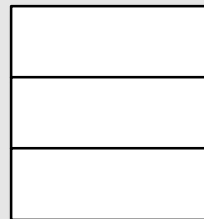
This implementation is very simple to program & understand since it's iterative

Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter
    (List<T> list,
     Function<T, T> op,
     ForkJoinPool fjPool) {
    fjPool.invoke
        (new RecursiveTask<List<T>>() {
            protected List<T> compute() {
                ...
            }
        })
    ...
}
```

ForkJoinPool

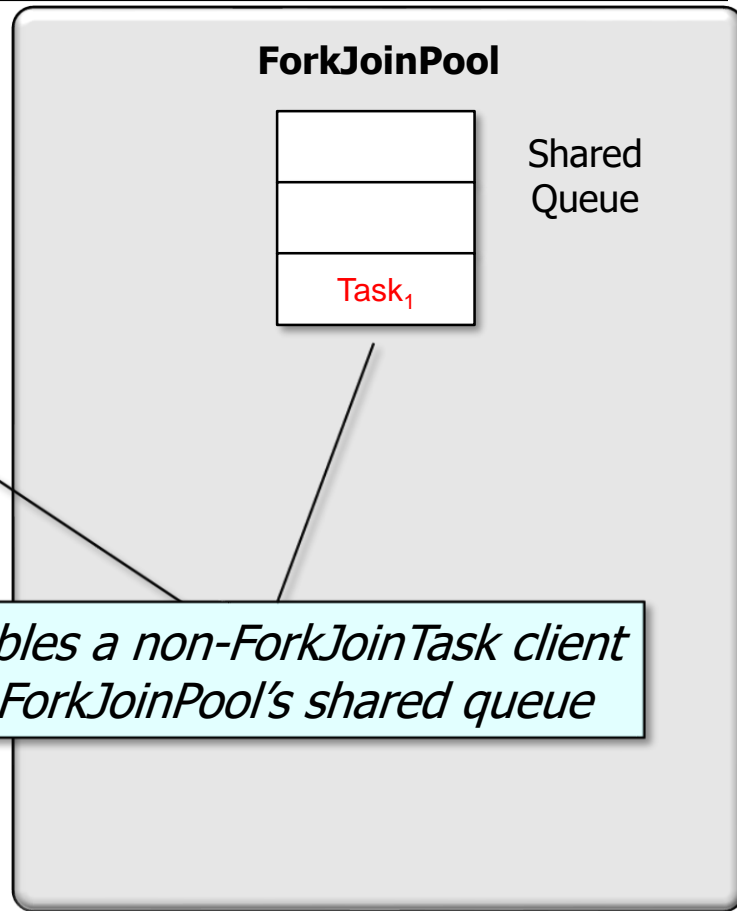


Shared
Queue

Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter
    (List<T> list,
     Function<T, T> op,
     ForkJoinPool fjPool) {
    fjPool.invoke
        (new RecursiveTask<List<T>>() {
            protected List<T> compute() {
                ...
            }
        })
    ...
}
```



The invoke() method enables a non-ForkJoinTask client to insert a task into the ForkJoinPool's shared queue

Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {
```

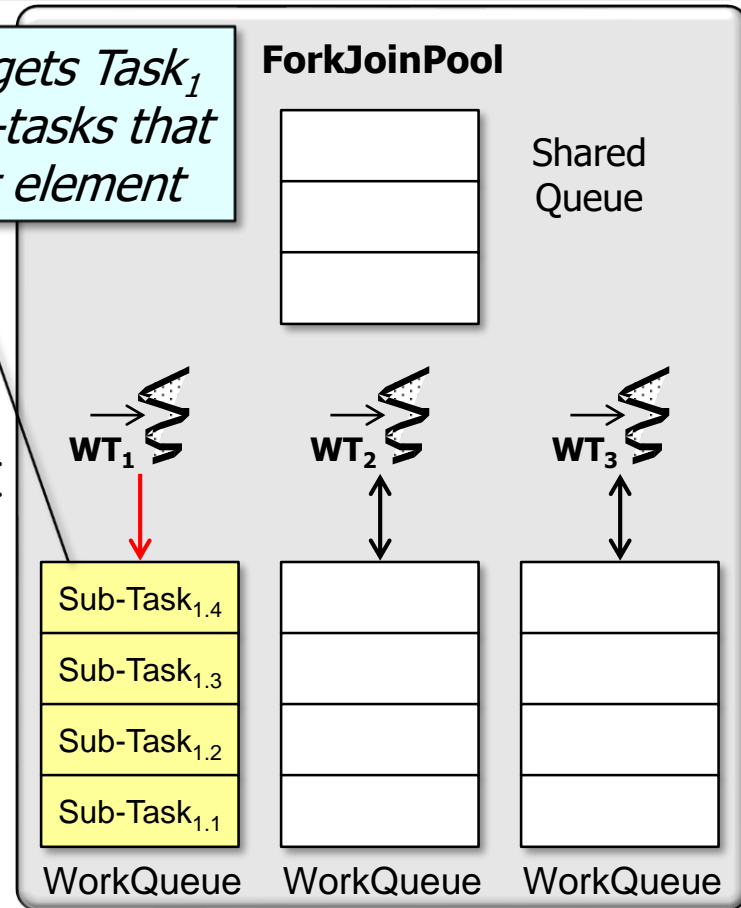
...

```
    for (T t : list)  
        forks.add(new RecursiveTask<T>() {  
            protected T compute()  
            { return op.apply(t); }  
        }.fork());
```

```
    for (ForkJoinTask<T> task : forks)  
        results.add(task.join());
```

...

Worker thread WT_1 gets $Task_1$ & creates n new sub-tasks that run 'op' on each list element



The highlighted code runs in the RecursiveTask's compute() method

Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {
```

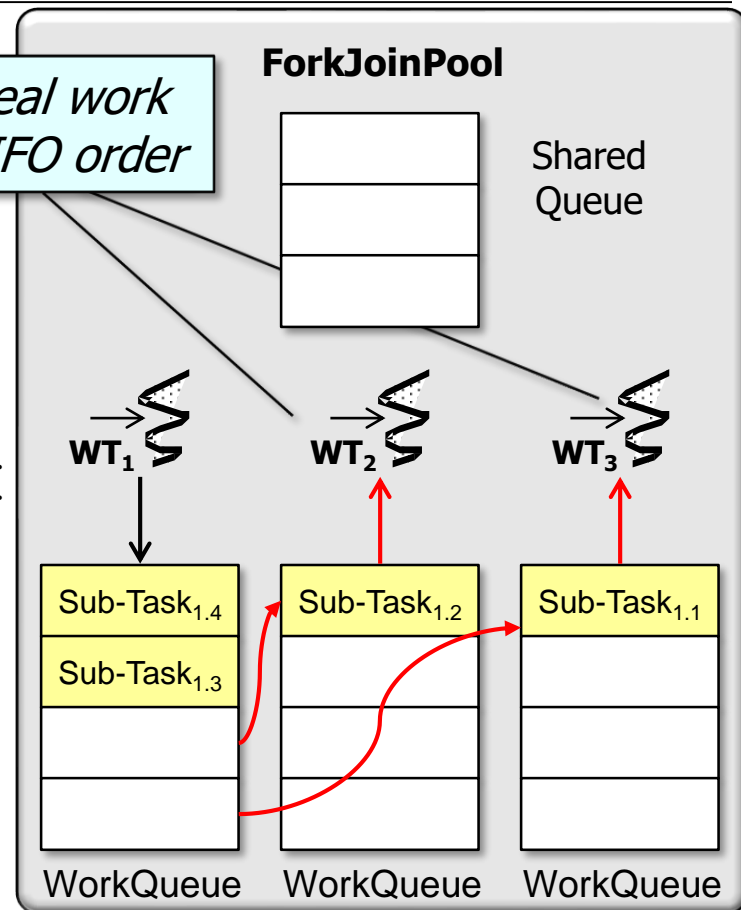
...

```
    for (T t : list)  
        forks.add(new RecursiveTask<T>() {  
            protected T compute()  
            { return op.apply(t); }  
        }.fork());
```

```
    for (ForkJoinTask<T> task : forks)  
        results.add(task.join());
```

...

*WT₂ & WT₃ steal work
from WT₁ in FIFO order*



“Work-stealing” overhead is high, but copying & method call overhead is low

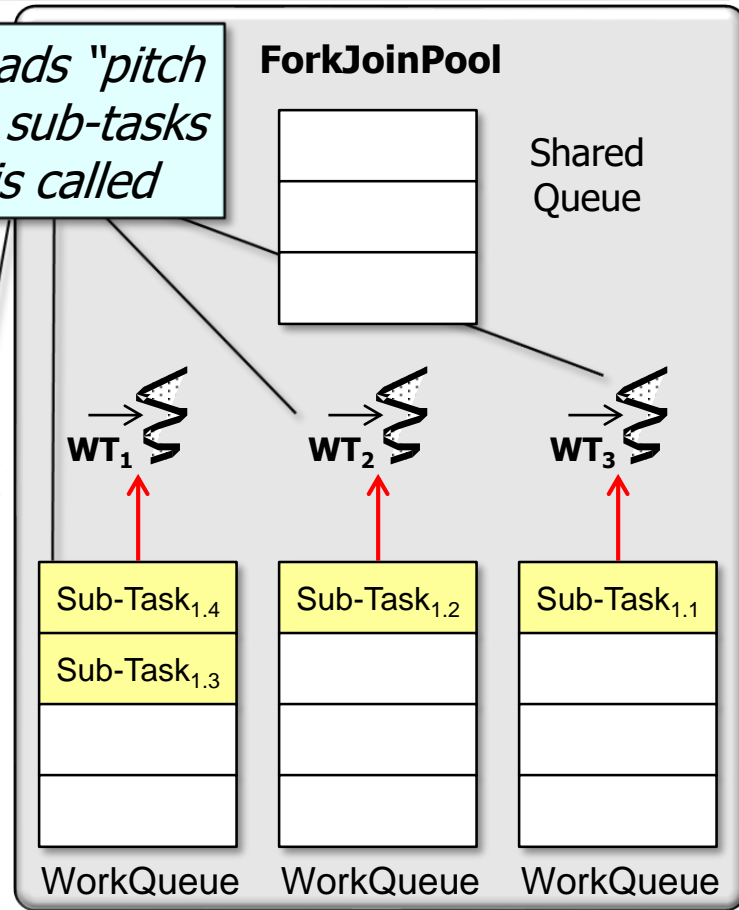
Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter
(List<T> list,
 Function<T, T> op,
 ForkJoinPool fjPool) {
    ...
    for (T t : list)
        forks.add(new RecursiveTask<T>() {
            protected T compute()
            { return op.apply(t); }
        }.fork());

    for (ForkJoinTask<T> task : forks)
        results.add(task.join());
    ...
}
```

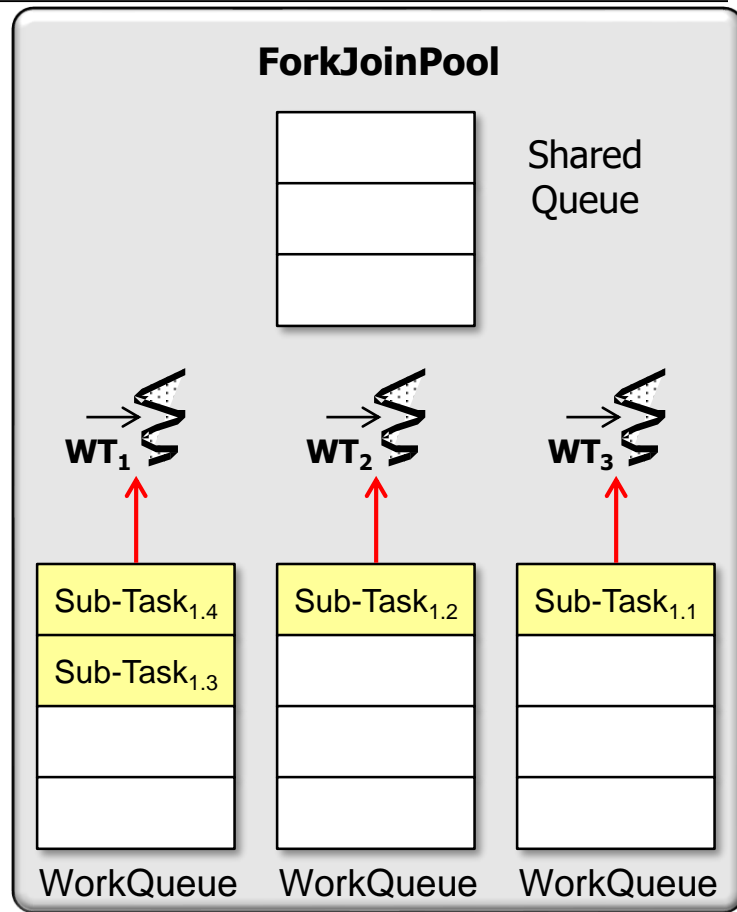
All worker threads "pitch in" to compute sub-tasks when join() is called



Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter
(List<T> list,
 Function<T, T> op,
 ForkJoinPool fjPool) {
...
for (T t : list)
    forks.add(new RecursiveTask<T>() {
        protected T compute()
        { return op.apply(t); }
    }.fork());
for (ForkJoinTask<T> task : forks)
    results.add(task.join());
...
}
```



"Collaborative Jiffy Lube" model of processing!

Implementing the applyAllIter() Method

- Visualizing applyAllIter()

```
<T> List<T> applyAllIter  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {
```

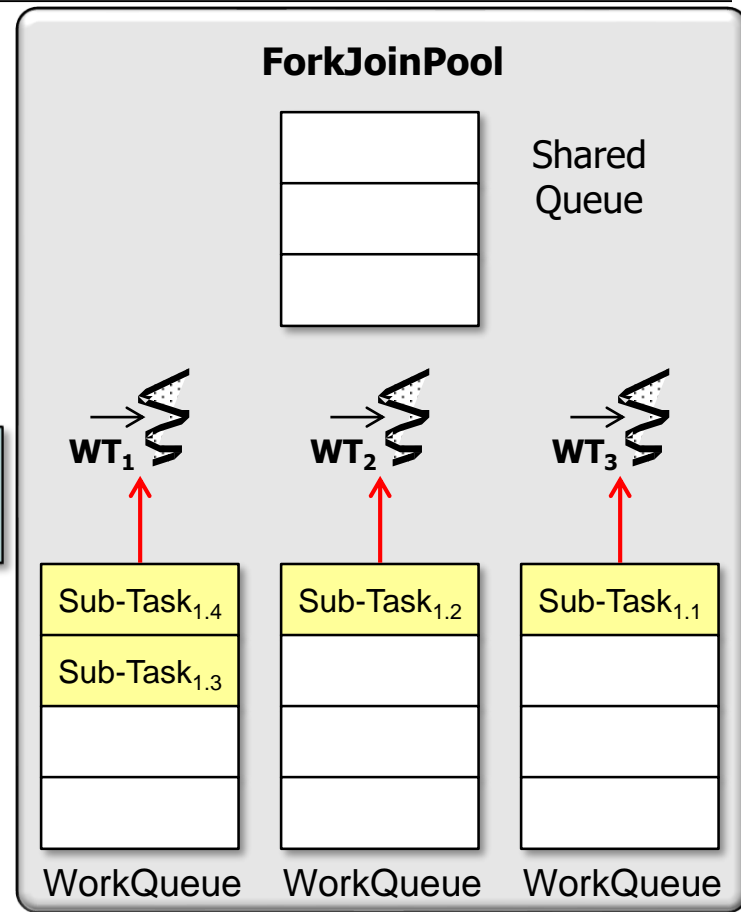
...

```
for (T t : list)  
    forks.add(new  
        protected T  
        { return op.apply(t); }  
    ).fork());
```

*This loop implements
"barrier synchronization"*

```
for (ForkJoinTask<T> task : forks)  
    results.add(task.join());
```

...



See [en.wikipedia.org/wiki/Barrier \(computer science\)](https://en.wikipedia.org/wiki/Barrier_(computer_science))

End of the Java Fork-Join Pool: Implementing `applyAllIter()`