

The Java Fork-Join Pool: Implementing `applyAllSplitIndex()`

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Apply the fork-join framework in practice
- Examine the `applyAllIter()` method
- Examine the `applyAllSplit()` method
- Examine the `applyAllSplitIndex()` method

```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {  
    ...
```

```
class SplitterTask  
extends RecursiveAction  
{ ... }
```

```
fjPool.invoke(new SplitterTask  
              (0, list.size()));
```

```
return Arrays.asList(results);  
}
```

Implementing the Method `applyAllSplitIndex()`

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) {  
    T[] results = (T[]) Array.newInstance(list.get(0).getClass(),  
                                           list.size());
```

```
    class SplitterTask extends RecursiveAction { ... }
```

```
    fjPool.invoke(new SplitterTask(0, list.size()));
```

```
    return Arrays.asList(results);
```

```
}
```

See [LiveLessons/blob/master/Java8/ex22/src/utils/ForkJoinUtils.java](https://livelessons.blob/master/Java8/ex22/src/utils/ForkJoinUtils.java)

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) {  
    T[] results = (T[]) Array.newInstance(list.get(0).getClass(),  
                                           list.size());
```

Create a new array to hold the results (yes, it's ugly...)

```
class SplitterTask extends RecursiveAction { ... }
```

```
fjPool.invoke(new SplitterTask(0, list.size()));
```

```
return Arrays.asList(results);
```

```
}
```

See stackoverflow.com/a/18137953

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) {  
    T[] results = (T[]) Array.newInstance(list.get(0).getClass(),  
                                           list.size());
```

This task partitions list recursively & runs each half in a ForkJoinTask

```
class SplitterTask extends RecursiveAction { ... }
```

```
fjPool.invoke(new SplitterTask(0, list.size()));
```

```
return Arrays.asList(results);
```

```
}
```

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) {  
    T[] results = (T[]) Array.newInstance(list.get(0).getClass(),  
                                           list.size());
```

```
    class SplitterTask extends RecursiveAction { ... }
```

```
    fjPool.invoke(new SplitterTask(0, list.size()));
```

```
    return Arrays.asList(results);
```

```
}
```



Invoke a new SplitterTask in the fork-join pool & wait for results

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) {  
    T[] results = (T[]) Array.newInstance(list.get(0).getClass(),  
                                           list.size());
```

```
    class SplitterTask extends RecursiveAction { ... }
```

```
    fjPool.invoke(new SplitterTask(0, list.size()));
```

```
    return Arrays.asList(results);
```

```
}
```

*Create & return a list
from the array of results*

These conversions to & from list → array → list incur overhead

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndexEx(List<T> list,  
                                   Function<T, T> op,  
                                   ForkJoinPool fjPool,  
                                   T[] results) {
```

*An alternative—more elegant & more efficient—
approach passes the results array as a parameter*

```
class SplitterTask extends RecursiveAction {  
    /* same implementation as follows next.. */  
}  
  
fjPool.invoke(new SplitterTask(0, list.size()));  
}
```

This approach is cleaner & faster since it has no conversion overhead

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...
```

```
class SplitterTask extends RecursiveAction {
```

```
    private int mLo;
```

```
    private int mHi;
```

```
    private SplitterTask(int lo, int hi) {
```

```
        mLo = lo;
```

```
        mHi = hi;
```

```
    }
```

```
    ...
```

```
} ...
```



*This task partitions list recursively
& runs each half in a ForkJoinTask*

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...
```

```
class SplitterTask extends RecursiveAction {
```

```
    private int mLo;
```

```
    private int mHi;
```

It uses indices to avoid the overhead of copy sub-lists

```
    private SplitterTask(int lo, int hi) {
```

```
        mLo = lo;
```

```
        mHi = hi;
```

```
    }
```

```
    ...
```

```
} ...
```

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...
```

```
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute();  
            lt.join();  
        } ...  
    }  
}
```

*Recursively perform the
computations in parallel*

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...
```

```
class SplitterTask extends RecursiveAction {
```

```
    protected void compute() {
```

```
        int mid = (mLo + mHi) >>> 1;
```

Find mid-point in current range

```
        if (mLo == mid)
```

```
            results[mLo] = op.apply(list.get(mLo));
```

```
        else {
```

```
            ForkJoinTask<Void> lt =
```

```
                new SplitterTask(mLo, mLo = mid).fork();
```

```
            compute();
```

```
            lt.join();
```

```
        } ...
```



This code recursively & evenly partitions the list

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...  
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute();  
            lt.join();  
        } ...  
    }  
}
```

*Apply op if there's
just one element*

This if statement handles the base case for the recursion

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...  
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute();  
            lt.join();  
        } ...  
    }  
}
```

Create a new task to handle the left-hand side of the list & fork it

This implementation uses recursive decomposition to disperse tasks to worker threads

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...  
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute() ;  
            lt.join();  
        } ...  
    }  
}
```

*Compute the right-hand side
in parallel with left-hand side*

`compute()` runs in the same task as its "parent" to minimize context switching

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...  
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute();  
            lt.join();  
        } ...  
    }  
}
```

*Join with left-hand side (this
is a synchronization point)*

Implementing the Method `applyAllSplitIndex()`

- Apply an 'op' to all list items by recursively splitting via fork-join method calls

```
<T> List<T> applyAllSplitIndex(List<T> list,  
                                Function<T, T> op,  
                                ForkJoinPool fjPool) { ...  
class SplitterTask extends RecursiveAction {  
    protected void compute() {  
        int mid = (mLo + mHi) >>> 1;  
        if (mLo == mid)  
            results[mLo] = op.apply(list.get(mLo));  
        else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo, mLo = mid).fork();  
            compute();  
            lt.join();  
        } ...  
    }  
}
```

This implementation is also harder to program & understand since it's recursive

Implementing the Method `applyAllSplitIndex()`

- Visualizing `applyAllSplitIndex()`

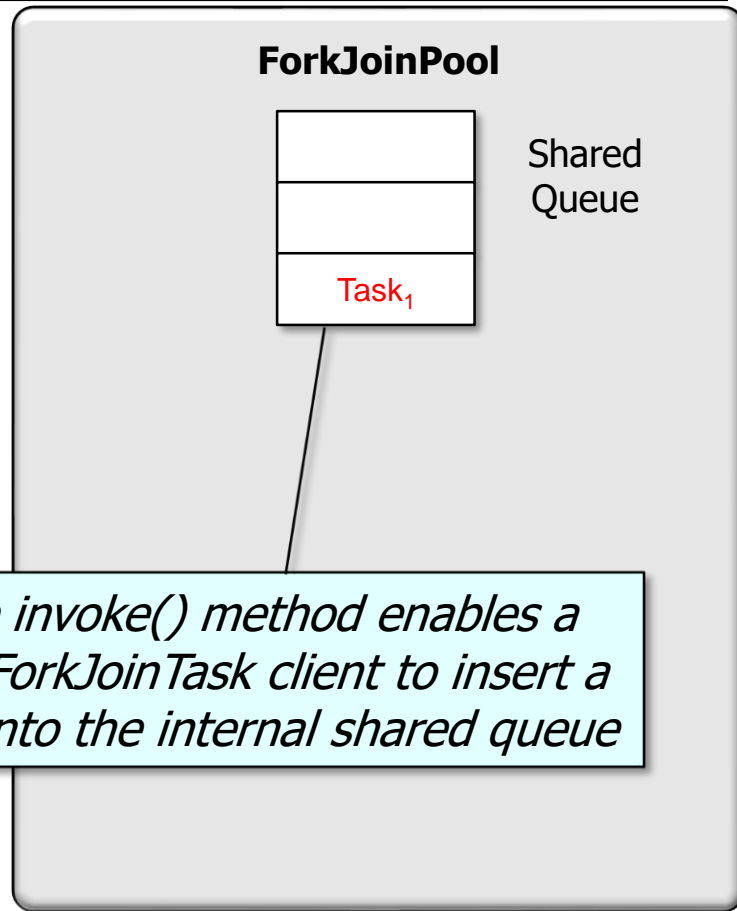
```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {
```

...

fjPool

```
    .invoke(new SplitterTask  
            (0,  
             list.size())));
```

...



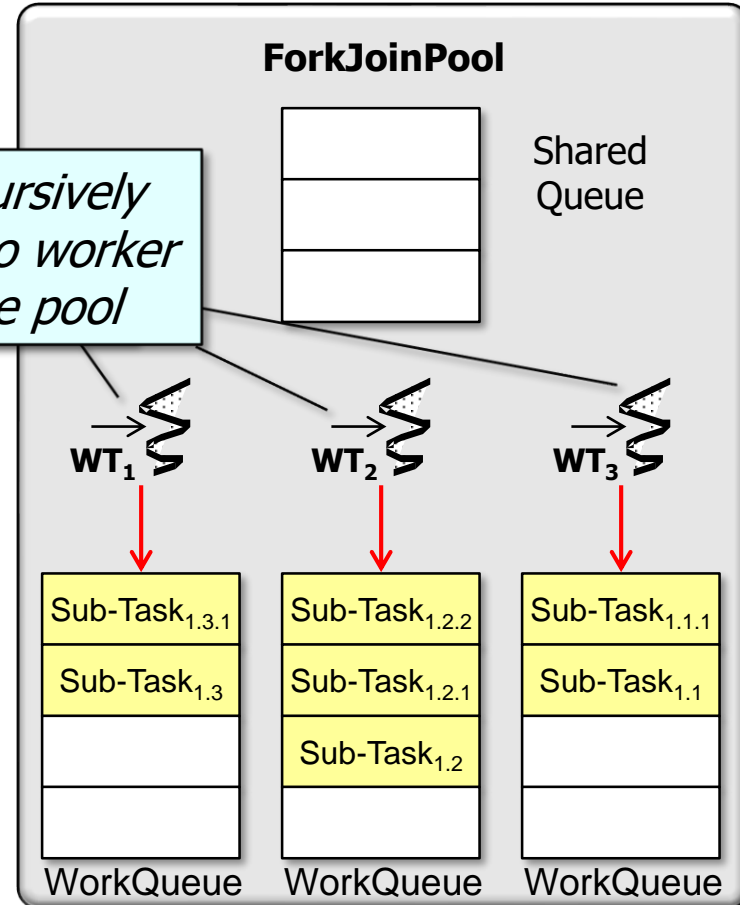
The `invoke()` method enables a non-ForkJoinTask client to insert a task into the internal shared queue

Implementing the Method `applyAllSplitIndex()`

- Visualizing `applyAllSplitIndex()`

```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T>  
     ForkJoinPool fj)  
class SplitterTask ... {  
    protected void compute() {  
        ... else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo,  
                                mLo = mid)  
                .fork();  
            compute();  
            lt.join();  
        }  
    }  
}
```

*Sub-tasks recursively
decompose onto worker
threads in the pool*



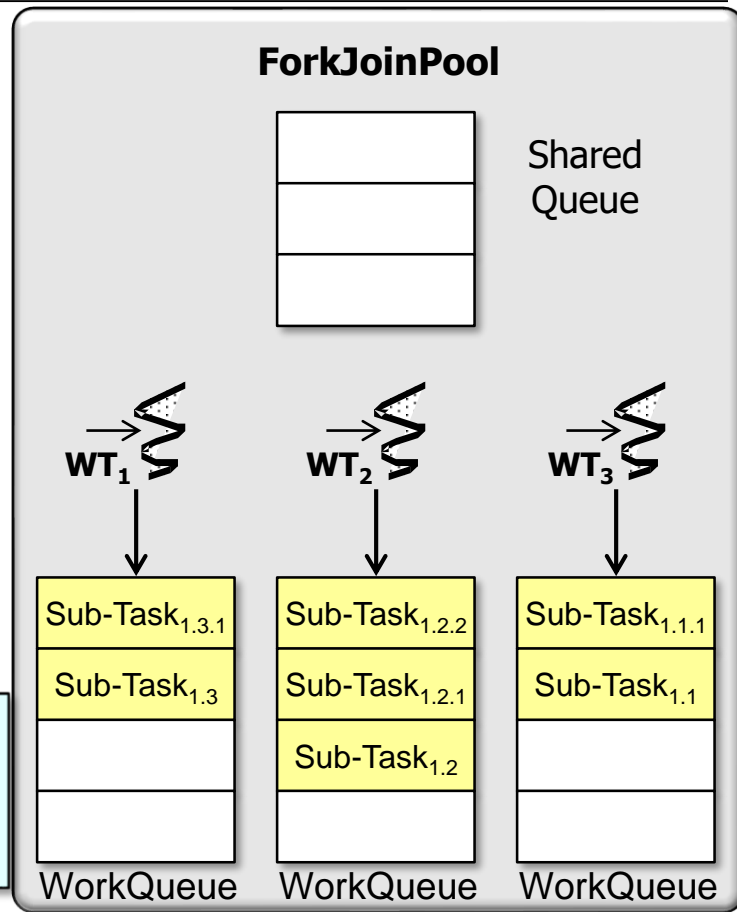
“Work-stealing” & copying overhead is low, but method call overhead is high

Implementing the Method `applyAllSplitIndex()`

- Visualizing `applyAllSplitIndex()`

```
<T> List<T> applyAllSplitIndex
(List<T> list,
 Function<T, T> op,
 ForkJoinPool fjPool) { ...
class SplitterTask ... {
    protected void compute() {
        ... else {
            ForkJoinTask<Void> lt =
                new SplitterTask(mLo,
                                mLo = mid)
                .fork();
            compute();
            lt.join();
        }
    }
}
```

*The fork()'d sub-task
& compute() sub-task
can run in parallel*



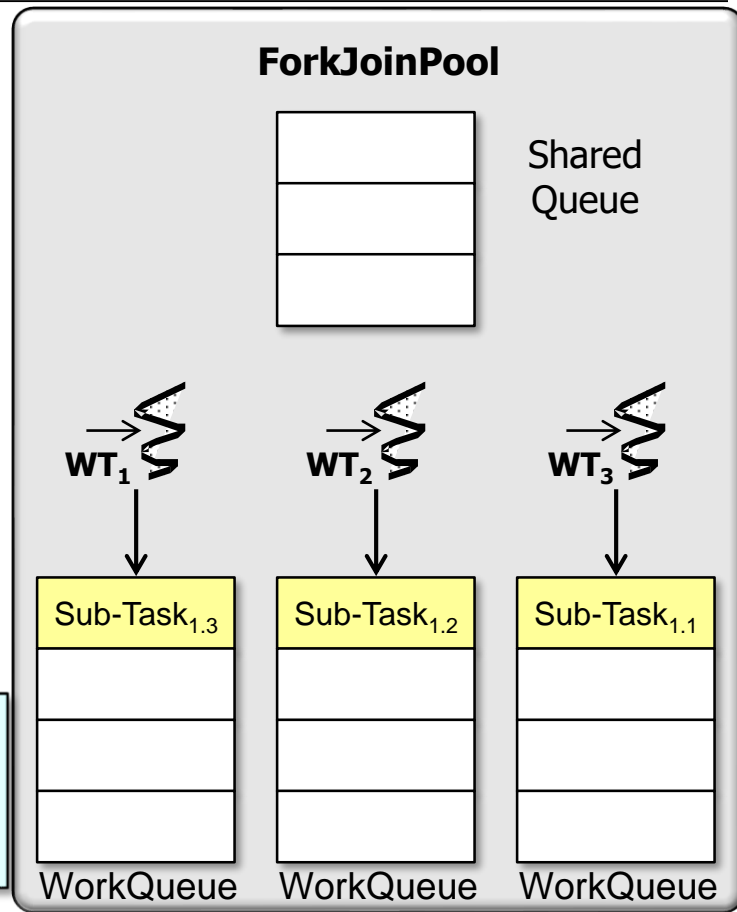
Implementing the Method `applyAllSplitIndex()`

- Visualizing `applyAllSplitIndex()`

```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) { ...
```

```
class SplitterTask ... {  
    protected void compute() {  
        ... else {  
            ForkJoinTask<Void> lt =  
                new SplitterTask(mLo,  
                                mLo = mid)  
                .fork();  
            compute();  
            lt.join();  
        }  
    }
```

*join() returns no value
& just serves as a
synchronization point*



There is a “balanced tree” of `join()` calls

End of the Java Fork-Join Pool: Implementing `applyAllSplitIndex()`