

The Java Fork-Join Pool: Evaluating the Example Applications

Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt



Professor of Computer Science

**Institute for Software
Integrated Systems**

**Vanderbilt University
Nashville, Tennessee, USA**



Learning Objectives in this Part of the Lesson

- Apply the fork-join framework in practice
- Examine the `applyAllIter()` method
- Examine the `applyAllSplit()` method
- Examine the `applyAllSplitIndex()` method
- Evaluate the example applications of the Fork-Join Pool framework



Evaluating the Example Applications

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons



Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand



```
<T> List<T> applyAllIter
    (List<T> list,
     Function<T, T> op,
     ForkJoinPool fjPool) {
    ...
    for (T t : list)
        forks.add
            (new RecursiveTask<T>() {
                protected T compute()
                { return op.apply(t); }
            }).fork();

    for (ForkJoinTask<T> task : forks)
        results.add(task.join());
    ...
}
```

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - but it incurs more work-stealing



Starting ForkJoinTest

applyAllIter() steal count = 31

applyAllSplitIndex() steal count = 16

applyAllSplit() steal count = 21

applyAllSplitIndexEx() steal count = 21

[1] Printing 4 results from fastest to slowest

testApplyAllSplitIndexEx() executed in 4575 ms

testApplyAllSplitIndex() executed in 5145 ms

testApplyAllSplit() executed in 5172 ms

testApplyAllIter() executed in 5599 ms

[1] Finishing ForkJoinTest

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - but it incurs more work-stealing
 - which lowers performance



Starting ForkJoinTest

applyAllIter() steal count = 31

applyAllSplitIndex() steal count = 16

applyAllSplit() steal count = 21

applyAllSplitIndexEx() steal count = 21

[1] Printing 4 results from fastest to slowest

testApplyAllSplitIndexEx() executed in 4575 ms

testApplyAllSplitIndex() executed in 5145 ms

testApplyAllSplit() executed in 5172 ms

testApplyAllIter() executed in 5599 ms

[1] Finishing ForkJoinTest

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”

Starting ForkJoinTest

applyAllIter() steal count = 31

applyAllSplitIndex() steal count = 16

applyAllSplit() steal count = 21

applyAllSplitIndexEx() steal count = 21

[1] Printing 4 results from fastest to slowest

testApplyAllSplitIndexEx() executed in 4575 ms

testApplyAllSplitIndex() executed in 5145 ms


testApplyAllSplit() executed in 5172 ms

testApplyAllIter() executed in 5599 ms

[1] Finishing ForkJoinTest

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - which improves performance



Starting ForkJoinTest

applyAllIter() steal count = 31

applyAllSplitIndex() steal count = 16

applyAllSplit() steal count = 21

applyAllSplitIndexEx() steal count = 21

[1] Printing 4 results from fastest to slowest

testApplyAllSplitIndexEx() executed in 4575 ms

testApplyAllSplitIndex() executed in 5145 ms

testApplyAllSplit() executed in 5172 ms

testApplyAllIter() executed in 5599 ms

[1] Finishing ForkJoinTest

There are also other factors (e.g., less data copying) that improve performance

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative `fork()/join()` is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - which improves performance
 - but is more complicated to program



```
class SplitterTask extends
    RecursiveTask<List<T>> {
protected List<T> compute() {
    ...
    int mid = mList.size() / 2;
    ForkJoinTask<List<T>> lt =
        new SplitterTask(mList.subList
            (0, mid)).fork();
    mList = mList
        .subList(mid, mList.size());
    List<T> rightResult = compute();
    List<T> leftResult = lt.join();
    leftResult.addAll(rightResult);
    return leftResult;
} ...
```

Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative `fork()/join()` is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - which improves performance
 - but is more complicated to program
 - & also does more “work” wrt method calls, etc.



Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - RecursiveAction’s overhead is lower than RecursiveTask’s

Starting ForkJoinTest

applyAllIter() steal count = 31

applyAllSplitIndex() steal count = 16

applyAllSplit() steal count = 21

applyAllSplitIndexEx() steal count = 21

[1] Printing 4 results from fastest to slowest

testApplyAllSplitIndexEx() executed in 4575 ms

testApplyAllSplitIndex() executed in 5145 ms

testApplyAllSplit() executed in 5172 ms

testApplyAllIter() executed in 5599 ms

[1] Finishing ForkJoinTest



Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - RecursiveAction’s overhead is lower than RecursiveTask’s
 - But RecursiveAction is also more idiosyncratic

```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {  
    T[] results = (T[]) Array  
        .newInstance  
        (list.get(0).getClass(),  
         list.size());  
    ...
```



Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - RecursiveAction’s overhead is lower than RecursiveTask’s
 - But RecursiveAction is also more idiosyncratic
 - Especially for generics

```
<T> List<T> applyAllSplitIndex  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool) {  
    T[] results = (T[]) Array  
        .newInstance  
            (list.get(0).getClass(),  
             list.size());  
    ...  
}
```



Evaluating the Example Applications

- Each Java fork-join programming model has pros & cons, e.g.
 - Iterative fork()/join() is simple to program/understand
 - Recursive decomposition incurs fewer “steals”
 - RecursiveAction’s overhead is lower than RecursiveTask’s
 - But RecursiveAction is also more idiosyncratic
 - Especially for generics
 - Changing the API can help!

```
<T> List<T> applyAllSplitIndexEx  
    (List<T> list,  
     Function<T, T> op,  
     ForkJoinPool fjPool,  
     T[] results) {
```

...



End of the Java Fork-Join Pool: Evaluating the Example Applications